



# 第2章 信息的表示与处理

100076202: 计算机系统导论

浮点数

Floating Point



任课教师:

宿红毅 张艳 黎有琦 颜珂

Carnegie  
Mellon  
University

原作者:

Randal E. Bryant and David R. O'Hallaron



# 议题: 浮点数 Floating Point

- 背景: 二进制小数 Background: Fractional binary numbers
- IEEE浮点标准: 定义 IEEE floating point standard: Definition
- 示例和属性 Example and properties
- 舍入、加法和乘法 Rounding, addition, multiplication
- C语言中的浮点数 Floating point in C
- 小结 Summary



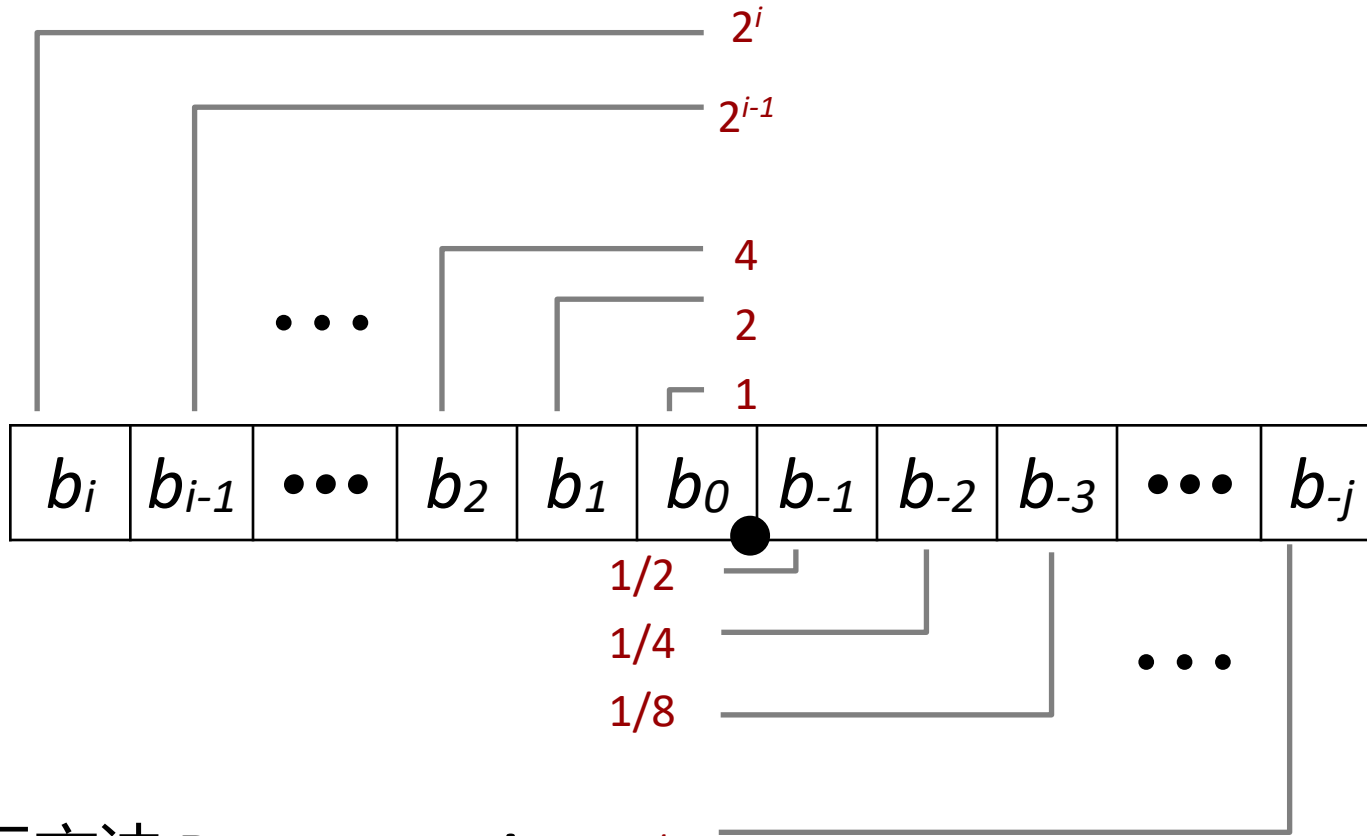
# 二进制小数

## Fractional binary numbers

- What is  $1011.101_2$ ?



# 二进制小数 Fractional Binary Numbers



## ■ 表示方法 Representation $2^{-j}$

- “小数点”右边的位代表2的整数次幂分之一 Bits to right of “binary point” represent fractional powers of 2

- 代表有理数 Represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$



# 二进制小数：示例

## Fractional Binary Numbers: Examples

### ■ 值 Value                      表示 Representation

5  $3/4 = 23/4$                        $101.11_2$

2  $7/8 = 23/8$                        $10.111_2$

1  $7/16 = 23/16$                        $1.0111_2$

### ■ 观察 Observations

- 通过右移来除以2（无符号数） Divide by 2 by shifting right (unsigned)
- 通过左移乘以2 Multiply by 2 by shifting left
- 数字形式  $0.111111..._2$  是刚好低于1.0的数 Numbers of form  $0.111111..._2$  are just below 1.0
  - $1/2 + 1/4 + 1/8 + ... + 1/2^i + ... \rightarrow 1.0$
  - 使用记法为  $1.0 - \epsilon$  Use notation  $1.0 - \epsilon$

# 可表示的数 Representable Numbers



## ■ 限制#1 Limitation #1

- 仅可以精确地表示 $x/2^k$ 形式的数 Can only exactly represent numbers of the form  $x/2^k$ 
  - 其它有理数有重复的比特位表示 Other rational numbers have repeating bit representations
- 值 Value      表示 Representation
  - $1/3$        $0.0101010101 [01] \dots_2$
  - $1/5$        $0.001100110011 [0011] \dots_2$
  - $1/10$        $0.0001100110011 [0011] \dots_2$

## ■ 限制#2 Limitation #2 ---定点数

- 在 $w$ 比特位中仅有一个二进制小数点设置 Just one setting of binary point within the  $w$  bits
  - 有限的数值范围（非常小的值？非常大的值？） Limited range of numbers (very small values? very large?)



# 议题: 浮点数 Floating Point

- 背景: 二进制小数 Background: Fractional binary numbers
- **IEEE浮点数标准: 定义 IEEE floating point standard: Definition**
- 示例和属性 Example and properties
- 舍入、加法和乘法 Rounding, addition, multiplication
- C语言中的浮点数 Floating point in C
- 小结 Summary

# IEEE浮点数 IEEE Floating Point



## ■ IEEE 754标准 IEEE Standard 754

- 1985年制定作为浮点运算的统一标准 Established in 1985 as uniform standard for floating point arithmetic
  - 在此之前，有很多异质的格式 Before that, many idiosyncratic formats
- 得到所有主流CPU的支持 Supported by all major CPUs
- 由Kahan为Intel处理器设计（获得1989年图灵奖） Designed by W. Kahan for Intel processors (Turing Award 1989)

## ■ 由数值问题所驱动 Driven by numerical concerns

- 非常好的标准用于舍入、上溢和下溢 Nice standards for rounding, overflow, underflow
- 在硬件上很难快速运算 Hard to make fast in hardware
  - 数值分析师在定义标准时比硬件设计师更占主导地位 Numerical analysts predominated over hardware designers in defining standard





# 浮点表示 Floating Point Representation

## ■ 浮点数形式 Numerical Form:

$$(-1)^s M 2^E$$

Example:

$$15213_{10} = (-1)^0 \times 1.1101101101101_2 \times 2^{13}$$

- 符号位 $s$ 确定数值是负还是正 **Sign bit  $s$**  determines whether number is negative or positive
- 尾数 $M$ 是范围在 $[1.0, 2.0)$ 之间的普通小数 **Significand  $M$**  normally a fractional value in range  $[1.0, 2.0)$ .
- 阶码 $E$ 是给浮点数指定 $2$ 的 $E$ 次幂权重 **Exponent  $E$**  weights value by power of two

## ■ 编码 Encoding

- 最高位是符号位 $s$  MSB  $s$  is sign bit  $s$
- **exp**字段编码 $E$  (但不等于 $E$ ) **exp** field encodes  $E$  (but is not equal to  $E$ )
- **frac**字段编码 $M$  (但不等于 $M$ ) **frac** field encodes  $M$  (but is not equal to  $M$ )



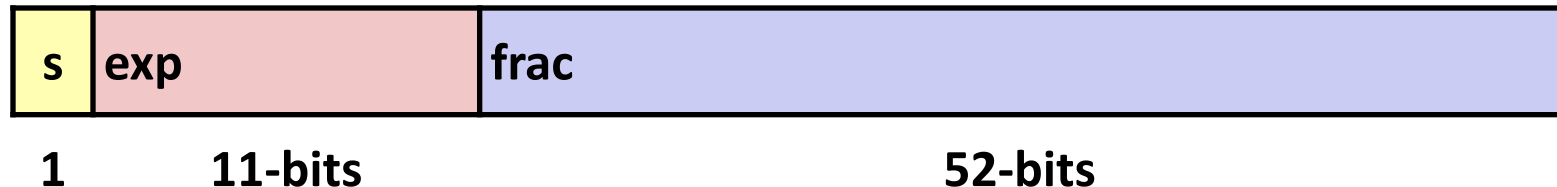


# 精度选项 Precision options

## ■ 单精度浮点数 Single precision: 32 bits



## ■ 双精度浮点数 Double precision: 64 bits



## ■ 扩展精度 (仅Intel) Extended precision: 80 bits (Intel only)



# “规格化”值

## “Normalized” Values

$$v = (-1)^s M 2^E$$



- 当阶码非全零和全一时 When:  $\text{exp} \neq 000\dots 0$  and  $\text{exp} \neq 111\dots 1$
- 阶码编码为一个有偏置值的有符号数:  $E = \text{Exp} - \text{Bias}$   
Exponent coded as a *biased* value:  $E = \text{Exp} - \text{Bias}$ 
  - $\text{Exp}$ :  $\text{exp}$ 字段的无符号值  $\text{Exp}$ : unsigned value of exp field
  - 偏置  $\text{Bias} = 2^{k-1} - 1$  其中  $k$  是阶码位的位数  $\text{Bias} = 2^{k-1} - 1$ , where  $k$  is number of exponent bits
    - 单精度: 127 Single precision: 127 ( $\text{Exp}$ : 1...254,  $E$ : -126...127)
    - 双精度: 1023 Double precision: 1023 ( $\text{Exp}$ : 1...2046,  $E$ : -1022...1023)
- 尾数编码为带隐含的一个前导1 Significand coded with implied leading 1:  $M = 1.\text{xxx}\dots\text{x}_2$ 
  - $\text{frac}$ 字段的比特位  $\text{xxx}\dots\text{x}$ : bits of frac field
  - 当  $\text{frac}$  全零时值最小 Minimum when  $\text{frac} = 000\dots 0$  ( $M = 1.0$ )
  - 当  $\text{frac}$  全一时值最大 Maximum when  $\text{frac} = 111\dots 1$  ( $M = 2.0 - \epsilon$ )

# 规格化编码示例



## Normalized Encoding Example

■ 值：单精度浮点数 Value: float  $F = 15213.0$ ;

$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \times 2^{13} \end{aligned}$$

$$\begin{aligned} v &= (-1)^s M 2^E \\ E &= \text{Exp} - \text{Bias} \end{aligned}$$

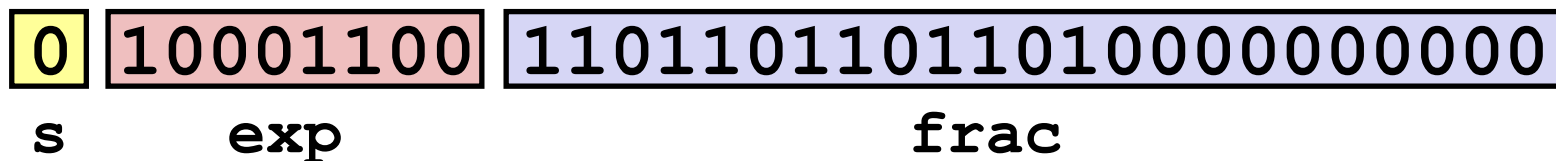
■ 尾数 Significand

$$\begin{aligned} M &= 1.\underline{1101101101101}_2 \\ \text{frac} &= \underline{11011011011010000000000}_2 \end{aligned}$$

■ 阶码 Exponent

$$\begin{aligned} E &= 13 \\ \text{Bias} &= 127 \\ \text{Exp} &= 140 = 10001100_2 \end{aligned}$$

■ Result:





# 非规格化值 Denormalized Values

- 条件:  $\text{exp}$  为全零 Condition:  $\text{exp} = 000\dots 0$

$$v = (-1)^s M 2^E$$
$$E = 1 - \text{Bias}$$

- 阶码值: Exponent value:  $E = 1 - \text{Bias}$  (instead of  $E = 0 - \text{Bias}$ )
- 尾数编码为隐含的一个前导零 Significand coded with implied leading 0:  $M = 0.\text{xxx}\dots\text{x}_2$ 
  - $\text{frac}$  字段比特位  $\text{xxx}\dots\text{x}$ : bits of  $\text{frac}$
- 情况 Cases
  - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$ 
    - 代表0值 Represents zero value
    - 注意区别值:  $+0$  和  $-0$  (为何?) Note distinct values:  $+0$  and  $-0$  (why?)
  - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$ 
    - 最接近0.0的数值 Numbers closest to 0.0
    - 平均分布的 Equispaced



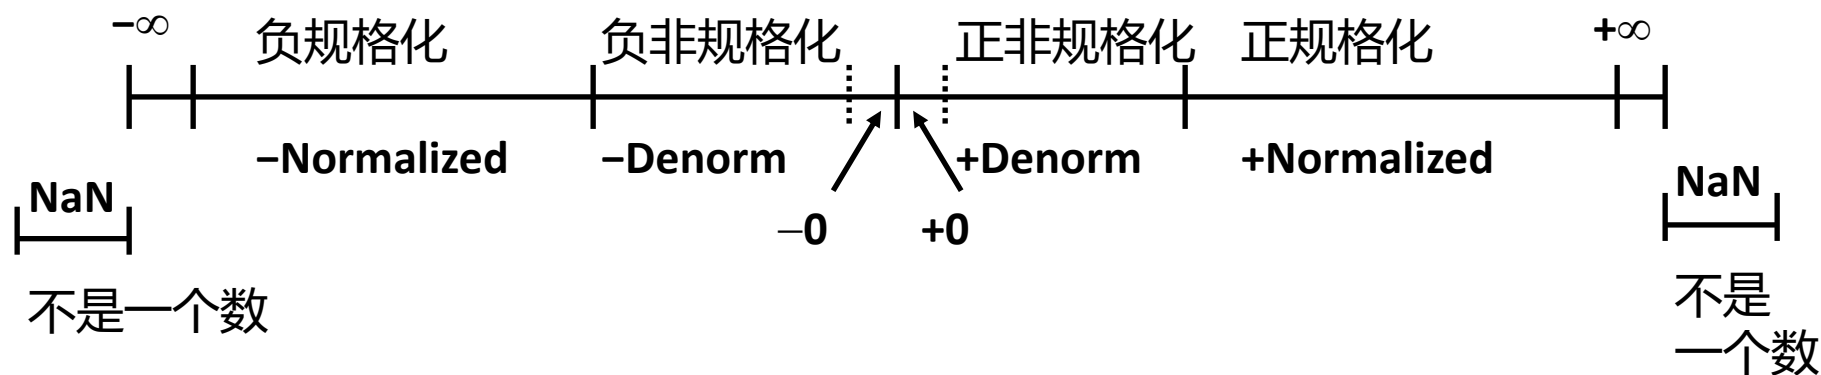
# 特殊值 Special Values

- 条件:  $\text{exp}$  为全一 Condition:  $\text{exp} = 111\dots 1$
- 情况 Case:  $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$ 
  - 代表值无穷大 Represents value  $\infty$  (infinity)
  - 溢出运算 Operation that overflows
  - 正负均如此 Both positive and negative
  - 例如 E.g.,  $1.0/0.0 = -1.0/-0.0 = +\infty$ ,  $1.0/-0.0 = -\infty$
- 情况 Case:  $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$ 
  - 不是一个数 Not-a-Number (NaN)
  - 代表无法确定数值的情况 Represents case when no numeric value can be determined
  - 例如 E.g.,  $\text{sqrt}(-1)$ ,  $\infty - \infty$ ,  $\infty \times 0$



# 可视化：浮点编码

## Visualization: Floating Point Encodings



# C语言中的浮点数示例

float: 0xC0A00000

binary: \_\_\_\_\_



1

8-bits

23-bits

**E =**

**S =**

**M =**

**$v = (-1)^S M 2^E =$**

$$v = (-1)^S M 2^E$$
$$E = \text{exp} - \text{Bias}$$

$$\text{Bias} = 2^{k-1} - 1 = 127$$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111





# C语言中的浮点数示例#1

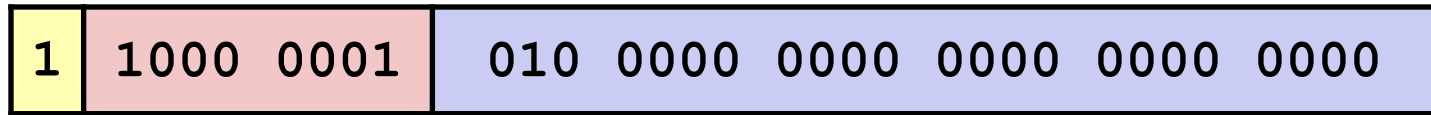
float: 0xC0A00000

$$v = (-1)^s M 2^E$$

$$E = \text{exp} - \text{Bias}$$

$$\text{Bias} = 2^{k-1} - 1 = 127$$

binary: 1100 0000 1010 0000 0000 0000 0000 0000



1

8-bits

23-bits

$$E = \text{exp} - \text{Bias} = 129 - 127 = 2 \text{ (decimal)}$$

$$S = 1 \rightarrow \text{负数}$$

$$M = 1.010 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000$$

$$= 1 + 1/4 = 1.25$$

$$v = (-1)^s M 2^E = (-1)^1 * 1.25 * 2^2 = -5$$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

# C语言中的浮点数示例#2



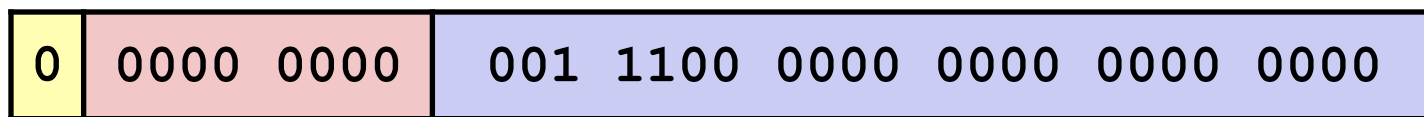
$$v = (-1)^s M 2^E$$

$$E = 1 - \text{Bias}$$

float: 0x001C0000

$$\text{Bias} = 2^{k-1} - 1 = 127$$

binary: 0000 0000 0001 1100 0000 0000 0000 0000



1

8-bits

23-bits

$$E = 1 - \text{Bias} = 1 - 127 = -126 \text{ (decimal)}$$

$$S = 0 \rightarrow \text{正数}$$

$$M = 0.001\ 1100\ 0000\ 0000\ 0000\ 0000$$

$$= 1/8 + 1/16 + 1/32 = 7/32 = 7 * 2^{-5}$$

$$v = (-1)^s M 2^E = (-1)^0 * 7 * 2^{-5} * 2^{-126} = 7 * 2^{-131}$$

$$\approx 2.571393892 \times 10^{-39}$$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111



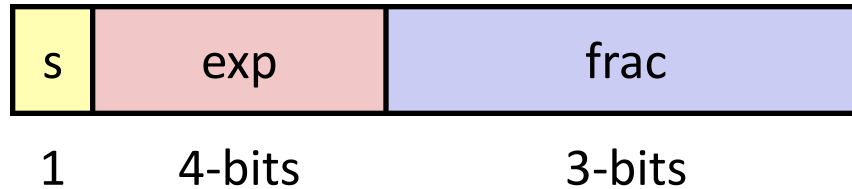
# 议题: 浮点数 Floating Point

- 背景: 二进制小数 Background: Fractional binary numbers
- IEEE浮点数标准: 定义 IEEE floating point standard: Definition
- 示例和属性 **Example and properties**
- 舍入、加法和乘法 Rounding, addition, multiplication
- C语言中的浮点数 Floating point in C
- 小结 Summary



# 微小的浮点数示例

## Tiny Floating Point Example



### ■ 8位浮点表示 8-bit Floating Point Representation

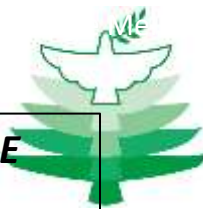
- 符号位是在最高有效位 the sign bit is in the most significant bit
- 随后四位是阶码，偏置为7 the next four bits are the exponent, with a bias of 7
- 最后的三位是尾数 the last three bits are the **frac**

### ■ 与IEEE格式同样的通用形式 Same general form as IEEE Format

- 规格化、非规格化 normalized, denormalized
- 0、NaN和无穷大的表示 representation of 0, NaN, infinity

# 动态范围 (仅正数)

## Dynamic Range (Positive Only)



s	exp	frac	E	Value
0	0000	000	-6	0
0	0000	001	-6	$1/8 * 1/64 = 1/512$
0	0000	010	-6	$2/8 * 1/64 = 2/512$
...				
0	0000	110	-6	$6/8 * 1/64 = 6/512$
0	0000	111	-6	$7/8 * 1/64 = 7/512$
0	0001	000	-6	$8/8 * 1/64 = 8/512$
0	0001	001	-6	$9/8 * 1/64 = 9/512$
...				
0	0110	110	-1	$14/8 * 1/2 = 14/16$
0	0110	111	-1	$15/8 * 1/2 = 15/16$
0	0111	000	0	$8/8 * 1 = 1$
0	0111	001	0	$9/8 * 1 = 9/8$
0	0111	010	0	$10/8 * 1 = 10/8$
...				
0	1110	110	7	$14/8 * 128 = 224$
0	1110	111	7	$15/8 * 128 = 240$
0	1111	000	n/a	inf

$$v = (-1)^s M 2^E$$

*n:  $E = Exp - Bias$*

*d:  $E = 1 - Bias$*

最接近0 closest to zero

最大非规格化数 largest denorm

最小规格化数 smallest norm

最接近1以下 closest to 1 below

最接近1以上 closest to 1 above

最大规格化数 largest norm

无穷大 inf

非规格化数  
Denormalized  
numbers

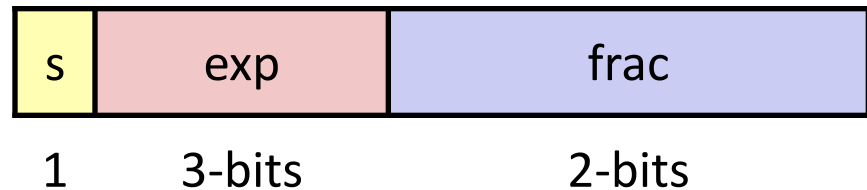
规格化数  
Normalized  
numbers



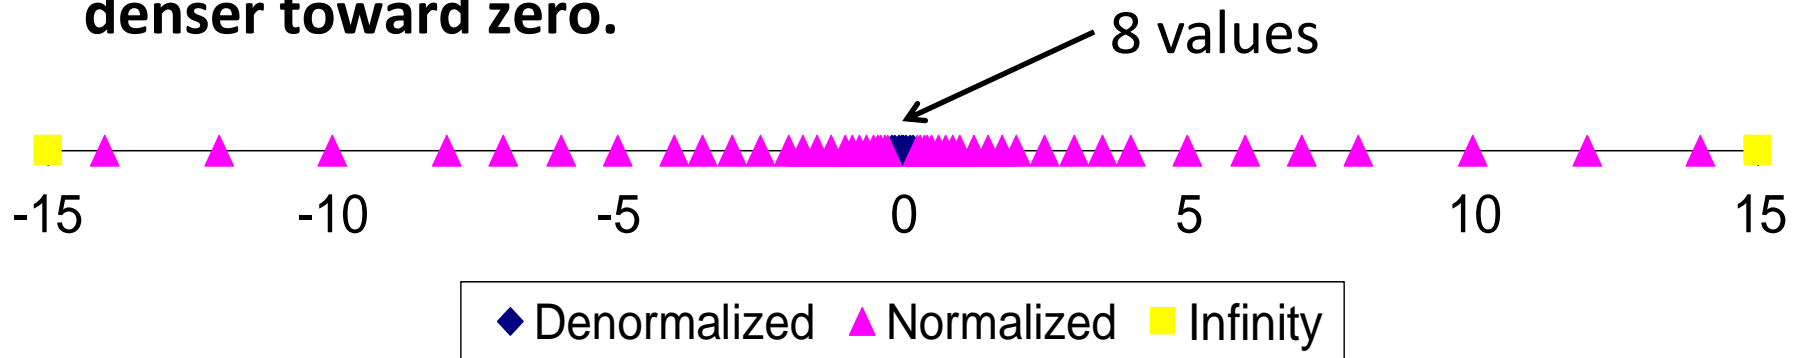
# 值的分布 Distribution of Values

## ■ 6位类IEEE格式 6-bit IEEE-like format

- 3位阶码  $e = 3$  exponent bits
- 2位尾数  $f = 2$  fraction bits
- 偏置是3 Bias is  $2^{3-1}-1 = 3$



## ■ 注意到越接近零分布越密集 Notice how the distribution gets denser toward zero.



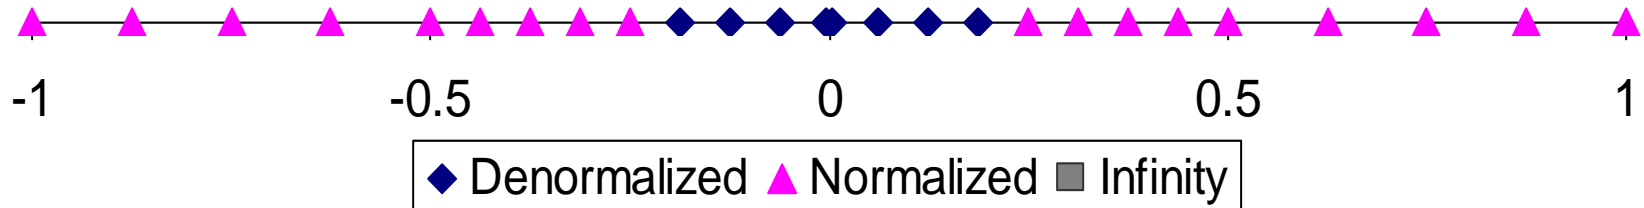
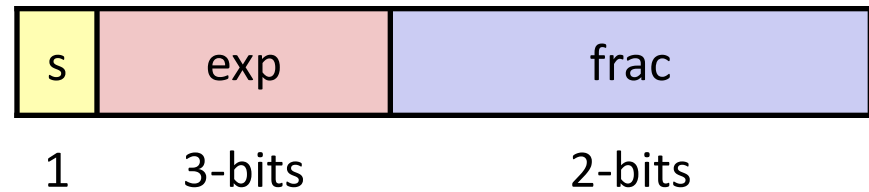


# 值的分布（近处观察）

## Distribution of Values (close-up view)

### ■ 6位类IEEE格式 6-bit IEEE-like format

- 3位阶码  $e = 3$  exponent bits
- 2位尾数  $f = 2$  fraction bits
- 偏置量是3 Bias is 3



# 有趣的数值 Interesting Numbers

{single, double}



<i>Description</i>	<i>exp</i>	<i>frac</i>	<i>Numeric Value</i>
■ Zero	00...00	00...00	0.0
■ Smallest Pos. Denorm.	00...00	00...01	$2^{-\{23,52\}} \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> <li>■ Single <math>\approx 1.4 \times 10^{-45}</math></li> <li>■ Double <math>\approx 4.9 \times 10^{-324}</math></li> </ul>			
■ Largest Denormalized	00...00	11...11	$(1.0 - \epsilon) \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> <li>■ Single <math>\approx 1.18 \times 10^{-38}</math></li> <li>■ Double <math>\approx 2.2 \times 10^{-308}</math></li> </ul>			
■ Smallest Pos. Normalized	00...01	00...00	$1.0 \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> <li>■ Just larger than largest denormalized</li> </ul>			
■ One	01...11	00...00	1.0
■ Largest Normalized	11...10	11...11	$(2.0 - \epsilon) \times 2^{\{127,1023\}}$
<ul style="list-style-type: none"> <li>■ Single <math>\approx 3.4 \times 10^{38}</math></li> <li>■ Double <math>\approx 1.8 \times 10^{308}</math></li> </ul>			





# IEEE编码的特殊属性

## Special Properties of the IEEE Encoding

- 浮点数的零和整数的零相同 **FP Zero Same as Integer Zero**
  - 所有比特位为0 All bits = 0
- 几乎可以使用无符号整数比较 **Can (Almost) Use Unsigned Integer Comparison**
  - 必须首先比较符号位 Must first compare sign bits
  - 必须考虑-0=0 Must consider  $-0 = 0$
  - 不是一个数NaN的问题 NaNs problematic
    - 比任何其它值都大 Will be greater than any other values
    - 还应该比较什么? What should comparison yield?
  - 否则都没有问题 Otherwise OK
    - 非规格化和规格化 Denorm vs. normalized
    - 规格化和无穷大 Normalized vs. infinity



# 议题: 浮点数 Floating Point

- 背景: 二进制小数 Background: Fractional binary numbers
- IEEE浮点数标准: 定义 IEEE floating point standard: Definition
- 示例和属性 Example and properties
- 舍入、加法和乘法 Rounding, addition, multiplication
- C语言中的浮点数 Floating point in C
- 小结 Summary



# 舍入 Rounding

## ■ 舍入模式（用美元来说明） Rounding Modes (illustrate with \$ rounding)

■		\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
■	向零 Towards zero	\$1	\$1	\$1	\$2	-\$1
■	向下 Round down ( $-\infty$ )	\$1	\$1	\$1	\$2	-\$2
■	向上 Round up ( $+\infty$ )	\$2	\$2	\$2	\$3	-\$1
■	最近的偶数（默认）	\$1	\$2	\$2	\$2	-\$2
■	Nearest Even (default)					

\*将数字向上或者向下舍入，使得结果的最低有效位是偶数

# 近处观察向偶数舍入



## Closer Look at Round-To-Even

### ■ 默认的舍入模式 Default Rounding Mode

- 没有深入汇编级很难理解任何其它类型的舍入 Hard to get any other kind without dropping into assembly
- 所有其它的舍入都有统计偏差 All others are statistically biased
  - 一组正数的和将始终高估或低估 Sum of set of positive numbers will consistently be over- or under- estimated

### ■ 适用于其它小数位/比特位位置 Applying to Other Decimal Places / Bit Positions

- 当正好位于两个可能值中间时 When exactly halfway between two possible values
  - 舍入以便最低位是偶数 Round so that least significant digit is even
- 例如舍入到最近的百分位 E.g., round to nearest hundredth

7.8949999	7.89	(低于中间值 Less than half way)
7.8950001	7.90	(高于中间值 Greater than half way)
7.8950000	7.90	(中间值-向上舍入 Half way—round up)
7.8850000	7.88	(中间值-向下舍入 Half way—round down)



# 舍入二进制数 Rounding Binary Numbers

## ■ 二进制小数 Binary Fractional Numbers

- 当最低位是0时为偶数 “Even” when least significant bit is 0
- 当舍入位置右边=100...<sub>2</sub>时为“中间值” “Half way” when bits to right of rounding position = 100...<sub>2</sub>

## ■ 举例 Examples

- 舍入到最接近1/4 (小数点右边2位) Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
2 3/32	10.00 <b>011</b> <sub>2</sub>	10.00 <sub>2</sub>	(<1/2—down向下)	2
2 3/16	10.00 <b>110</b> <sub>2</sub>	10.01 <sub>2</sub>	(>1/2—up向上)	2 1/4
2 7/8	10.11 <b>100</b> <sub>2</sub>	11.00 <sub>2</sub>	( 1/2—up向上)	3
2 5/8	10.10 <b>100</b> <sub>2</sub>	10.10 <sub>2</sub>	( 1/2—down向下)	2 1/2

# 舍入 Rounding



监督位：结果的最低位

Guard bit: LSB of result

舍入位：删除的第一位

Round bit: 1<sup>st</sup> bit removed

1 . BBG**RXXX**

固着位：剩余位的或

Sticky bit: OR of remaining bits

## ■ 向上舍入条件 Round up conditions

- Round = 1, Sticky = 1  $\rightarrow$   $> 0.5$
- Guard = 1, Round = 1, Sticky = 0  $\rightarrow$  Round to even 向偶数舍入

<i>Value</i>	<i>Fraction</i>	<i>GRS</i>	<i>Incr?</i>	<i>Rounded</i>
128	1.000 <b>0000</b>	000	N	1.000
15	1.101 <b>0000</b>	100	N	1.101
17	1.000 <b>1000</b>	010	N	1.000
19	1.001 <b>1000</b>	110	Y	1.010
138	1.000 <b>1010</b>	011	Y	1.001
63	1.111 <b>1100</b>	111	Y	10.000



# 浮点运算：基本思想

## Floating Point Operations: Basic Idea

■  $x +_f y = \text{Round}(x + y)$

■  $x \times_f y = \text{Round}(x \times y)$

■ 基本思想 Basic idea

- 首先计算精确的结果 First **compute exact result**
- 使它适合需要的精度 Make it fit into desired precision
  - 如果阶码太大可能会溢出 Possibly overflow if exponent too large
  - 可能需要舍入才能适合尾数位数 Possibly **round to fit into frac**

# 浮点加法 Floating Point Addition



■  $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

■ 假设 Assume  $E1 > E2$

小数点对齐 Get binary points lined up

■ 精确的结果 Exact Result:  $(-1)^s M 2^E$

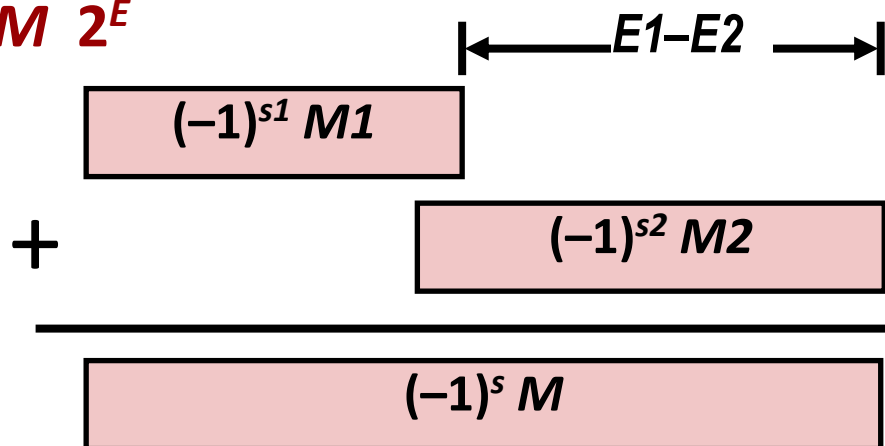
■ 符号位  $s$ , 尾数  $M$ :

■ Sign  $s$ , significand  $M$ :

■ 有符号数对齐并相加的结果

■ Result of signed align & add

■ 阶码  $E$  是  $E1$  Exponent  $E$ :  $E1$



## ■ 修正 Fixing

■ 如果  $M$  大于等于 2,  $M$  右移,  $E$  加一 If  $M \geq 2$ , shift  $M$  right, increment  $E$

■ 如果  $M$  小于 1,  $M$  左移  $k$  位,  $E$  减  $k$  if  $M < 1$ , shift  $M$  left  $k$  positions, decrement  $E$  by  $k$

■ 如果  $E$  超过范围则溢出 Overflow if  $E$  out of range

■ 舍入  $M$  到适合  $\text{frac}$  的精度 Round  $M$  to fit **frac** precision





# 浮点加法 Floating Point Addition

$$\begin{aligned} 1.010 * 2^2 + 1.110 * 2^3 &= (0.1010 + 1.1100) * 2^3 \\ &= 1\textcolor{red}{0}.0110 * 2^3 = 1.001\textcolor{red}{10} * 2^4 = 1.010 * 2^4 \end{aligned}$$



# 浮点加法的数学性质

## Mathematical Properties of FP Add

### ■ 相比于其它阿贝尔群 Compare to those of Abelian Group

- 加法封闭吗? Closed under addition? **Yes**
  - 但可能产生无穷大或NaN But may generate infinity or NaN
- 可交换吗? Commutative? **Yes**
- 可结合吗? Associative? **No**
  - 溢出和舍入不精确 Overflow and inexactness of rounding
  - $(3.14 + 1e10) - 1e10 = 0$ ,  $3.14 + (1e10 - 1e10) = 3.14$
- 0是加性恒等 (单位元) 的吗? 0 is additive identity? **Yes**
- 每个元素都有加法逆元吗? Every element has additive inverse?
  - 对, 除了无穷大和NaNs Yes, except for infinities & NaNs **Almost**

### ■ 单调性 Monotonicity

- $a \geq b \Rightarrow a+c \geq b+c$  **Almost**
  - 除了无穷大和NaNs Except for infinities & NaNs



# 浮点乘法 FP Multiplication

■  $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$

■ 精确的结果 Exact Result:  $(-1)^s M 2^E$

- 符号位s: 异或      Sign s:       $s1 \wedge s2$
- 尾数M: 相乘      Significand M:       $M1 \times M2$
- 阶码E: 相加      Exponent E:       $E1 + E2$

■ 修正 Fixing

- 如果M大于等于2, M右移, 阶码E加一 If  $M \geq 2$ , shift  $M$  right, increment  $E$
- 如果E超过范围, 溢出 If  $E$  out of range, overflow
- 舍入M到适合frac的精度 Round  $M$  to fit **frac** precision

■ 实现 Implementation

- 最繁琐的工作是尾数相乘 Biggest chore is multiplying significands

4 位尾数:  $1.010 \times 2^2 \times 1.110 \times 2^3 = 10.0011 \times 2^5$   
 $= 1.00011 \times 2^6 = 1.001 \times 2^6$

# 浮点乘法的数学性质



## Mathematical Properties of FP Mult

### ■ 相比于交换环 Compare to Commutative Ring

- 乘法封闭吗? Closed under multiplication? **Yes**
  - 但可能产生无穷大或NaN But may generate infinity or NaN
- 乘法可交换吗? Multiplication Commutative? **Yes**
- 乘法具有结合性吗? Multiplication is Associative? **No**
  - 可能溢出, 舍入不精确 Possibility of overflow, inexactness of rounding
  - Ex:  $(1e20 * 1e20) * 1e-20 = \text{inf}$ ,  $1e20 * (1e20 * 1e-20) = 1e20$
- 1是乘法恒等 (单位元) 的吗? 1 is multiplicative identity? **Yes**
- 乘法对加法是可分配的吗? Multiplication distributes over addition? **No**
  - 可能溢出, 舍入不精确 Possibility of overflow, inexactness of rounding
  - $1e20 * (1e20 - 1e20) = 0.0$ ,  $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$

### ■ 单调性 Monotonicity

- $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c$  **Almost**
  - 除了无穷大和NaNs Except for infinities & NaNs



# 议题: 浮点数 Floating Point

- 背景: 二进制小数 Background: Fractional binary numbers
- IEEE浮点数标准: 定义 IEEE floating point standard: Definition
- 示例和属性 Example and properties
- 舍入、加法和乘法 Rounding, addition, multiplication
- C语言中的浮点数 Floating point in C
- 小结 Summary

# C语言中的浮点数 Floating Point in C



## ■ C语言确保两个级别的浮点数 C Guarantees Two Levels

- **float**      single precision      单精度
- **double**      double precision      双精度

## ■ 转换/强制转换 Conversions/Casting

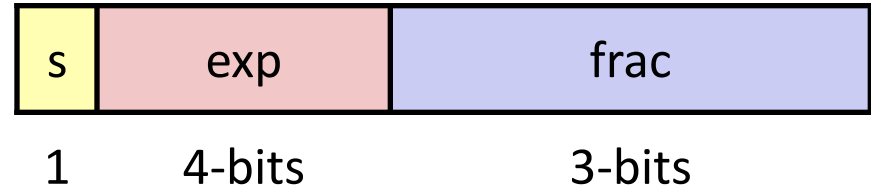
- 在int, float和double之间强制转换改变比特位表示 Casting between **int**, **float**, and **double** changes bit representation
- **double/float** → **int**
  - 截断尾数部分 Truncates fractional part
  - 就像向零舍入 Like rounding toward zero
  - 当超过范围或NaN时没有定义：一般设置为TMin Not defined when out of range or NaN: Generally sets to TMin
- **int** → **double**
  - 精确转换，只要int字长小于等于53位 Exact conversion, as long as **int** has  $\leq 53$  bit word size
- **int** → **float**
  - 将按照舍入模式进行舍入 Will round according to rounding mode

# 创建浮点数 Creating Floating Point Number



## ■ 步骤 Steps

- 用前导1规格化尾数 Normalize to have leading 1
- 舍入以适合尾数位 Round to fit within fraction
- 后规格化以处理舍入的影响 Postnormalize to deal with effects of rounding



## ■ 案例研究 Case Study

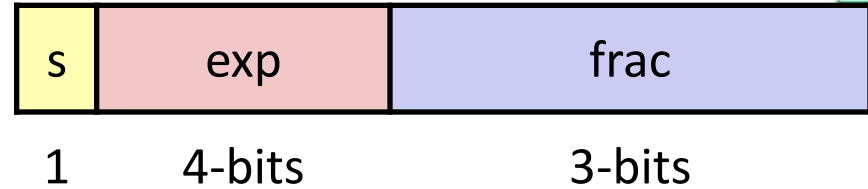
- 转换8位无符号数成微小浮点数格式 Convert 8-bit unsigned numbers to tiny floating point format

示例数值 Example Numbers

128	10000000
13	00001101
17	00010001
19	00010011
138	10001010
63	00111111



# 规格化 Normalize



## ■ 需求 Requirement

- 设置小数点以便数值格式为1.xxxxx Set binary point so that numbers of form 1.xxxxx
- 调整所有位得到前导1 Adjust all to have leading one
  - 随着尾数左移，阶码减一 Decrement exponent as shift left
  - **尾数右移，阶码+1**

<i>Value</i>	<i>Binary</i>	<i>Fraction</i>	<i>Exponent</i>
128	10000000	1.0000000	7
13	00001101	1.1010000	3
17	00010001	1.0001000	4
19	00010011	1.0011000	4
138	10001010	1.0001010	7
63	00111111	1.1111100	5





# 后规格化 Postnormalize

## ■ 问题 Issue

- 舍入可能导致溢出 Rounding may have caused overflow
- 通过右移一次进行处理同时阶码加一 Handle by shifting right once & incrementing exponent

<i>Value</i>	<i>Rounded</i>	<i>Exp</i>	<i>Adjusted</i>	<i>Result</i>
128	1.000	7		128
13	1.101	3		13
17	1.000	4		16
19	1.010	4		20
138	1.001	7		134
63	10.000	5	1.000/6	64



# 浮点数难题 Floating Point Puzzles

- 对于下面的每个C表达式 For each of the following C expressions, either: 完成其中一个工作
  - 对于所有的参数值解释其值为真 Argue that it is true for all argument values
  - 解释为何不为真 Explain why not true

```
int x = ...;  
float f = ...;  
double d = ...;
```

假设d和f都不是NaN  
Assume neither  
d nor f is NaN

- `x == (int)(float) x` ❌
- `x == (int)(double) x` ✅
- `f == (float)(double) f` ✅
- `d == (double)(float) d` ❌
- `f == -(-f);` ✅
- `2/3 == 2/3.0` ❌
- `d < 0.0`  $\Rightarrow$  `((d*2) < 0.0)` ✅
- `d > f`  $\Rightarrow$  `-f > -d` ✅
- `d * d >= 0.0` ✅
- `(d+f) - d == f` ❌



# 小结 Summary

- IEEE浮点数有明确的数学性质 IEEE Floating Point has clear mathematical properties
- 表示浮点数的形式为 Represents numbers of form  $M \times 2^E$
- 运算和实现是相互独立的 One can reason about operations independent of implementation
  - 好像采用完美的精度进行计算，然后进行舍入 As if computed with perfect precision and then rounded
- 与实数运算并不相同 Not the same as real arithmetic
  - 违背结合率和分配律 Violates associativity/distributivity
  - 对编译器和严谨数值应用程序员是一个挑战 Makes life difficult for compilers & serious numerical applications programmers

# 关于浮点数的灾难性影响 (I)



## Disastrous effects on floating Point (I)

- 浮点运算的不精确性 The imprecision of floating-point arithmetic
- 1991年2月25日，在第一次海湾战争期间，位于沙特阿拉伯 Dharan 的美国爱国者导弹连未能拦截来袭的伊拉克飞毛腿导弹。飞毛腿袭击了美国陆军营房并杀死了28名士兵。 On February 25, 1991, during the first Gulf War, an American **Patriot Missile** battery in Dharan, Saudi Arabia, failed to intercept an incoming Iraqi **Scud missile**. The Scud struck an American Army barracks and killed 28 soldiers.
- 爱国者系统包含一个内部时钟，以计数器的形式实现，每0.1 秒递增一次。为了确定以秒为单位的时间，程序会将此计数器的值乘以24比特位的小数，该小数是1/10的二进制近似值。 The Patriot system contains an internal clock, implemented as a counter that is incremented every 0.1 seconds. To determine the time in seconds, the program would multiply the value of this counter by a 24-bit quantity that was a fractional binary approximation to 1/10 .



# 关于浮点数的灾难性影响 (II)



## Disastrous effects on floating Point (II)

- 1996年6月4日，阿丽亚娜Ariane 5号火箭的处女航行 The maiden voyage of the Ariane 5 rocket, on June 4, 1996.
- 发射后仅37秒，火箭就偏离了飞行路径，解体并爆炸 Just 37 seconds after liftoff, the rocket veered off its flight path, broke up, and exploded.
- 在将64位浮点数转换为16位有符号整数的过程中发生了溢出 An overflow had occurred during the conversion of a 64-bit floating-point number to a 16-bit signed integer
- 溢出的值测量了火箭的水平速度 The value that overflowed measured the horizontal velocity of the rocket.
- 在Ariane 4中，水平速度永远不会溢出16位数字。Ariane 5只是以5倍的速度重用相同的软件 In the Ariane 4, the horizontal velocity would never overflow a 16-bit number. Ariane 5 just reuses the same software with 5 times higher velocity.

