



第12章 并发编程

同步：高级/Synchronization: Advanced

100076202： 计算机系统导论

任课教师：

宿红毅 张艳 黎有琦 颜珂

原作者：

Randal E. **Bryant** and David R. O'Hallaron



**Carnegie
Mellon
University**



议题 Today

- **回顾：信号量、互斥和生产者-消费者** Review: Semaphores, mutexes, producer-consumer
- **使用信号量调度共享资源** Using semaphores to schedule shared resources
 - **读者-写者问题** Readers-writers problem
- **其它并发问题** Other concurrency issues
 - **线程安全** Thread safety
 - **竞争** Races
 - **死锁** Deadlocks
 - **线程和信号处理之间交互** Interactions between threads and signal handling

提醒：信号量



Reminder: Semaphores

- **Semaphore:** non-negative global integer synchronization variable
- **Manipulated by P and V operations:**
 - $P(s)$: [**while** ($s == 0$) ; $s--$;]
 - Dutch for "Proberen" (test)
 - $V(s)$: [$s++$;]
 - Dutch for "Verhogen" (increment)
- **OS kernel guarantees that operations between brackets [] are executed atomically**
 - Only one P or V operation at a time can modify s .
 - When **while** loop in P terminates, only that P can decrement s
- **Semaphore invariant: ($s \geq 0$)**

回顾：使用信号量通过互斥保护共享资源



Review: Using semaphores to protect shared resources via mutual exclusion

■ 基本思想： Basic idea:

- 将一个唯一的信号量互斥锁 (*mutex*) (最初为1) 与每个共享变量 (或相关的共享变量集) 相关联 Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables)
- 用P(*mutex*)和V(*mutex*)操作包围对共享变量的每次访问 Surround each access to the shared variable(s) with *P(mutex)* and *V(mutex)* operations

```
mutex = 1
```

```
P(mutex)
```

```
cnt++
```

```
V(mutex)
```



回顾：使用锁进行互斥

Review: Using Lock for Mutual Exclusion

■ 基本思想： Basic idea:

- 互斥锁Mutex是只有值0（锁定）或1（解锁）的信号量的特殊情况
Mutex is special case of semaphore that only has value 0 (locked) or 1 (unlocked)
- *Lock(m)*: [**while** (m == 0) ; m=0;]
- *Unlock(m)*: [m=1]
- 比使用信号量快约2倍 ~2x faster than using semaphore for this purpose
- 而且，更清楚地表明程序员的意图 And, more clearly indicates programmer's intention

```
mutex = 1

lock(mutex)
cnt++
unlock(mutex)
```


关于示例的注释 **Note about Examples**



- **课程示例将使用信号量进行计数和互斥** **Lecture examples will use semaphores for both counting and mutual exclusion**
 - **代码比使用pthread_mutex短得多** **Code is much shorter than using pthread_mutex**

回顾：生产者-消费者问题

Review: Producer-Consumer Problem




■ 通用同步模式：Common synchronization pattern:

- 生产者等待空**槽位**，将项目插入缓冲区，并通知消费者 Producer waits for empty **slot**, inserts item in buffer, and notifies consumer
- 消费者等待**项目**，将其从缓冲区中删除，并通知生产者 Consumer waits for **item**, removes it from buffer, and notifies producer

回顾：生产者-消费者问题

Review: Producer-Consumer Problem



■ 示例 Examples

- 多媒体处理： Multimedia processing:
 - 生产者创建视频帧，消费者对其进行渲染 Producer creates video frames, consumer renders them
- 事件驱动的图形用户界面 Event-driven graphical user interfaces
 - 生产者检测鼠标点击、鼠标移动和键盘点击，并在缓冲区中插入相应的事件 Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - 消费者从缓冲区检索事件并绘制显示 Consumer retrieves events from buffer and paints the display

回顾：使用信号量协调共享资源的访问

Review: Using Semaphores to Coordinate Access to Shared Resources



- **基本思想：线程使用信号量操作通知另一个线程某些条件已变为真** **Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true**
 - **使用计数信号量来跟踪资源状态** Use counting semaphores to keep track of resource state.
 - **使用二元信号量通知其他线程** Use binary semaphores to notify other threads.

回顾：使用信号量协调共享资源的访问

Review: Using Semaphores to Coordinate Access to Shared Resources



■ 生产者-消费者问题 The Producer-Consumer Problem

- 对进程之间的交互活动进行协调，一个进程产生信息，另一个进程使用该信息 Mediating interactions between processes that generate information and that then make use of that information
- 用两个二元信号量实现单条目缓冲区 Single entry buffer implemented with two binary semaphores
 - 一个用于控制生产者的访问 One to control access by producer(s)
 - 一个用于控制消费者的访问 One to control access by consumer(s)
- 使用信号量+环形缓冲区实现N个条目缓冲区 N-entry implemented with semaphores + circular buffer




议题 Today

- 回顾：信号量、互斥和生产者-消费者 Review: Semaphores, mutexes, producer-consumer
- **使用信号量调度共享资源** Using semaphores to schedule shared resources
 - **读者-写者问题** Readers-writers problem
- **其它并发问题** Other concurrency issues
 - 线程安全 Thread safety
 - 竞争 Races
 - 死锁 Deadlocks
 - 线程和信号处理交互 Interactions between threads and signal handling

读者和写者问题

Readers-Writers Problem



■ 问题陈述: Problem statement:



- 读者线程仅读取对象 Reader threads only read the object
- 写者线程修改对象 (读/写访问) Writer threads modify the object (read/write access)
- 写者必须具有对对象的独占访问权限 Writers must have exclusive access to the object
- 无限数量的读者可以访问该对象 Unlimited number of readers can access the object

■ 在真实系统中频繁发生, 例如 Occurs frequently in real systems, e.g.,

- 在线航空预订系统 Online airline reservation system
- 多线程缓存Web代理 Multithreaded caching Web proxy

读者/写者示例

Readers/Writers Examples



读者和写者的变体 Variants of Readers-Writers



- 第一类读者-写者问题（有利于读者-读者优先） **First readers-writers problem (favors readers)**
 - 除非已授予写者使用对象的权限，否则不应让任何读者等待 No reader should be kept waiting unless a writer has already been granted permission to use the object.
 - 等待写者之后到达的读者比写者优先 A reader that arrives after a waiting writer gets priority over the writer.
- 第二类读者-写者问题（有利于写者-写者优先） **Second readers-writers problem (favors writers)**
 - 一旦写者准备好写入，它将尽快执行写入 Once a writer is ready to write, it performs its write as soon as possible
 - 在写者之后到达的读者必须等待，即使写者也要等待 A reader that arrives after a writer must wait, even if the writer is also waiting.
- 在这两种情况下都有可能出现饿死情况（线程无限期等待） **Starvation (where a thread waits indefinitely) is possible in both cases.**

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w);




        /* Writing here */

        V(&w);
    }
}
```

rw1.c

读者/写者示例

Readers/Writers Examples



第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

R1 → /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 1
W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        R2 → if (readcnt == 1) /* First in */
            P(&w);
            V(&mutex);

        R1 → /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 2
W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

R2
R1



写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 2
W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 1

W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        R3 → if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        R2 → P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);

        R1 → }
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 2

W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 1

W == 0

第一类读者-写者问题的解决方案

Solution to First Readers-Writers Problem



读者 Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

写者 Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

到达: Arrivals: R1 R2 W1 R3

Readcnt == 0

W == 1

R3 →

其它读者-写者版本

Other Versions of Readers-Writers



- **第一类解决方案的不足 Shortcoming of first solution**
 - 源源不断的读者将无限期地阻止写者 Continuous stream of readers will block writers indefinitely
- **第二个版本 Second version**
 - 一旦写者出现，就会阻止以后的读者访问 Once writer comes along, blocks access to later readers
 - 一系列写入可能会阻止所有读取 Series of writes could block all reads
- **先进先出实现 FIFO implementation**
 - 请参阅code目录中的rwqueue代码 See rwqueue code in code directory
 - 按顺序接收服务请求 Service requests in order received
 - 保存线程在先进先出队列中 Threads kept in FIFO
 - 每一个都有信号量，可以访问临界区 Each has semaphore that enables its access to critical section

第二类读者-写者问题解决方案

Solution to Second Readers-Writers Problem



```
int readcnt, writecnt;           // Initially 0
sem_t rmutex, wmutex, r, w;     // Initially 1
void reader(void)
{
    while (1) {
        P(&r);
        P(&rmutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&rmutex);
        V(&r)

        /* Reading happens here */

        P(&rmutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&rmutex);
    }
}
```

第二类读者-写者问题解决方案

Solution to Second Readers-Writers Problem




```
void writer(void)
{
    while (1) {
        P(&wmutex);
        writecnt++;
        if (writecnt == 1)
            P(&r);
        V(&wmutex);

        P(&w);
        /* Writing here */
        V(&w);

        P(&wmutex);
        writecnt--;
        if (writecnt == 0);
            V(&r);
        V(&wmutex);
    }
}
```

用先进先出队列管理读者/写者

Managing Readers/Writers with FIFO



思想 Idea

- 读/写请求插入先进先出队列 Read & Write requests are inserted into FIFO
- 请求在从队列删除时进行处理 Requests handled as remove from FIFO
 - 如果当前空闲或正在处理读取，则允许继续读取 Read allowed to proceed if currently idle or processing read
 - 仅允许在空闲时继续写入请求 Write allowed to proceed only when idle
- 请求完成后通知控制器 Requests inform controller when they have completed

公平 Fairness

- 保证最终会处理每个请求 Guarantee every request is eventually handled

读者写者先进先出实现

Readers Writers FIFO Implementation



- 完整代码见rwqueue.h和rwqueue.c Full code in rwqueue.{h,c}

```
/* Queue data structure */
typedef struct {
    sem_t mutex; // Mutual exclusion
    int reading_count; // Number of active readers
    int writing_count; // Number of active writers
    // FIFO queue implemented as linked list with tail
    rw_token_t *head;
    rw_token_t *tail;
} rw_queue_t;
```

```
/* Represents individual thread's position in queue */
typedef struct TOK {
    bool is_reader;
    sem_t enable; // Enables access
    struct TOK *next; // Allows chaining as linked list
} rw_token_t;
```

读者写者先进先出队列使用

Readers Writers FIFO Use



■ 在rwqueue-test.c程序中 In rwqueue-test.c

```
/* Get write access to data and write */
void iwriter(int *buf, int v)
{
    rw_token_t tok;
    rw_queue_request_write(&q, &tok);
    /* Critical section */
    *buf = v;
    /* End of Critical Section */
    rw_queue_release(&q);
}
```

```
/* Get read access to data and read */
int ireader(int *buf)
{
    rw_token_t tok;
    rw_queue_request_read(&q, &tok);
    /* Critical section */
    int v = *buf;
    /* End of Critical section */
    rw_queue_release(&q);
    return v;
}
```

读者/写者锁的库函数

Library Reader/Writer Lock



■ 数据类型 Data type `pthread_rwlock_t`

■ 操作 Operations

- 获得读锁 Acquire read lock

```
pthread_rwlock_rdlock(pthread_rwlock_t *rwlock)
```

- 获得写锁 Acquire write lock

```
pthread_rwlock_wrlock(pthread_rwlock_t *rwlock)
```

- 释放 (其中一个) 锁 Release (either) lock

```
pthread_rwlock_unlock(pthread_rwlock_t *rwlock)
```

■ 观察 Observation

- 必须正确使用库函数 Library must be used correctly!
 - 由程序员决定哪些需要读访问, 哪些需要写访问 Up to programmer to decide what requires read access and what requires write access



议题 Today

- 回顾：信号量、互斥和生产者-消费者 Review: Semaphores, mutexes, producer-consumer
- 使用信号量调度共享资源 Using semaphores to schedule shared resources
 - 读者-写者问题 Readers-writers problem
- 其它并发问题 **Other concurrency issues**
 - **竞争 Races**
 - 死锁 Deadlocks
 - 线程安全 Thread safety
 - 线程和信号处理之间交互 Interactions between threads and signal handling

一个担忧：竞争 One Worry: Races




- 当程序的正确性取决于一个线程在另一个线程到达点y之前到达点x时，就会发生**竞争** A **race** occurs when correctness of the program depends on one thread reaching point x before another thread reaches point y

```
/* a threaded program with a race */
int main(int argc, char** argv) {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++)
        Pthread_create(&tid[i], NULL, thread, &i);
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    return 0;
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

数据竞争 Data Race





消除竞争 Race Elimination

- **不要共享状态 Don't share state**
 - 例如, 使用malloc为每个线程生成单独的参数拷贝 E.g., use malloc to generate separate copy of argument for each thread
- **使用同步原语控制对共享状态的访问 Use synchronization primitives to control access to shared state**
 - 不同的共享状态可能使用不同的原语 Different shared state can use different primitives



议题 Today

- 回顾：信号量、互斥和生产者-消费者 Semaphores, mutexes, producer-consumer
- 使用信号量调度共享资源 Using semaphores to schedule shared resources
 - 生产者-消费者问题 Producer-consumer problem
- **其它并发问题 Other concurrency issues**
 - **竞争 Races**
 - **死锁 Deadlocks**
 - 线程安全 Thread safety
 - 线程和信号处理之间交互 Interactions between threads and signal handling



一个担忧：死锁 A Worry: Deadlock

- 定义：当且仅当一个进程正在等待一个永远不会为真的条件，那么它就会**死锁** Def: A process is **deadlocked** iff it is waiting for a condition that will never be true.
- 典型场景 Typical Scenario
 - 进程1和进程2需要两个资源（A和B）才能继续 Processes 1 and 2 needs two resources (A and B) to proceed
 - 进程1获取A，等待B Process 1 acquires A, waits for B
 - 进程2获取B，等待A Process 2 acquires B, waits for A
 - 两个进程都将永远等待！ Both will wait forever!



一个担忧：死锁 A Worry: Deadlock

- 定义：当且仅当一个进程正在等待一个永远不会为真的条件，那么它就会**死锁** Def: A process is **deadlocked** iff it is waiting for a condition that will never be true.
- 更全面的知识（超出了**213**课程的范围），死锁有四个要求
More fully (and beyond the scope of 213), a deadlock has four requirements
 - 互斥 Mutual exclusion
 - 循环等待 Circular waiting
 - 保持和等待 Hold and wait
 - 非抢占式 No pre-emption

信号量死锁 Deadlocking With Semaphores



```
int main(int argc, char** argv)
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1);  /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1);  /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    return 0;
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

Tid[0]:
P(s₀);
P(s₁);
cnt++;
V(s₀);
V(s₁);

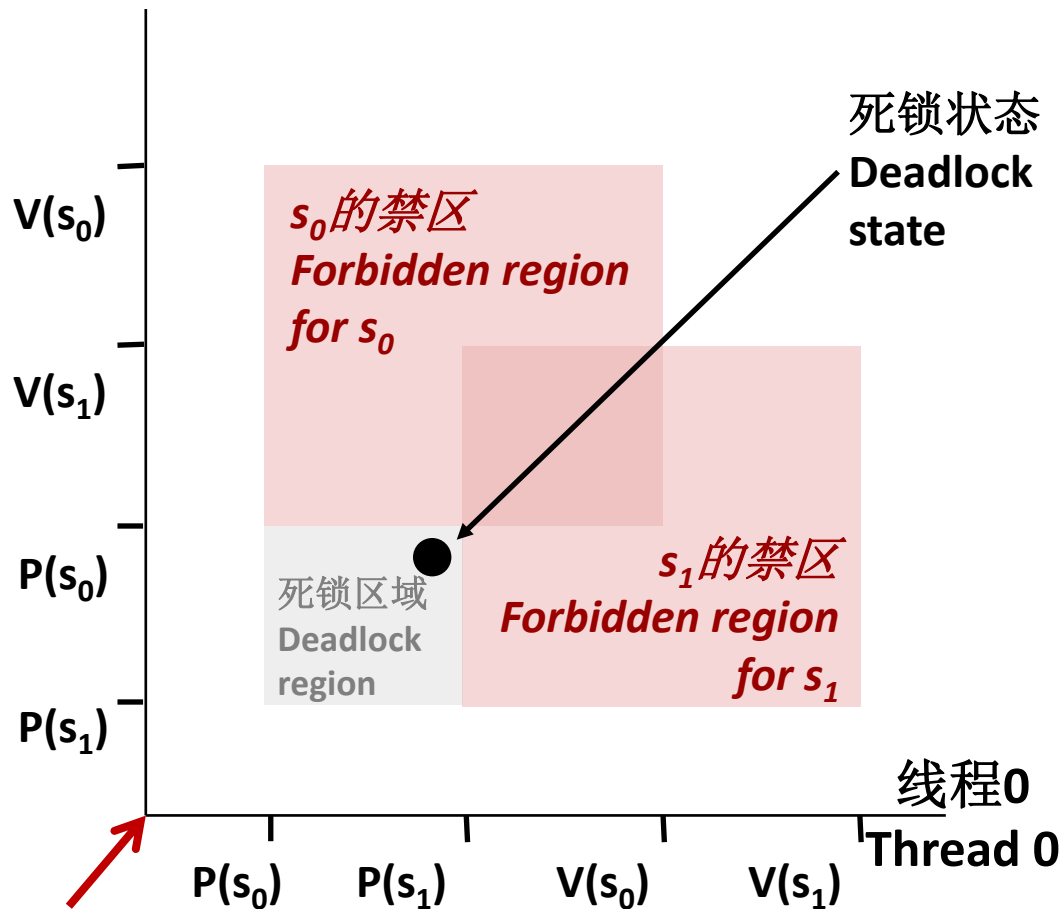
Tid[1]:
P(s₁);
P(s₀);
cnt++;
V(s₁);
V(s₀);

进度图中显示的死锁

Deadlock Visualized in Progress Graph



线程1 Thread 1



锁定引入了**死锁**的可能性：等待一个永远不会成真的条件 Locking introduces the potential for **deadlock**: waiting for a condition that will never be true

任何进入**死锁区域**的轨迹将最终到达**死锁状态**，等待 s_0 或 s_1 变为非零 Any trajectory that enters the **deadlock region** will eventually reach the **deadlock state**, waiting for either S_0 or S_1 to become nonzero

其他轨迹幸运地避开了死锁区域 Other trajectories luck out and skirt the deadlock region

不幸的事实：死锁往往是不确定的（竞争） Unfortunate fact: deadlock is often nondeterministic (race)

死锁 Deadlock



避免死锁

Avoiding Deadlock

以相同的顺序获取共享资源

Acquire shared resources in same order



```
int main(int argc, char** argv)
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1);  /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1);  /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    return 0;
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[0]); P(&mutex[1]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

Tid[0]:
P(s₀);
P(s₁);
cnt++;
V(s₀);
V(s₁);


Tid[1]:
P(s₀);
P(s₁);
cnt++;
V(s₁);
V(s₀);

在进度图中避免死锁

Avoided Deadlock in Progress Graph



线程1 Thread 1



轨迹无法卡住 No way for trajectory to get stuck

进程以相同的顺序获取锁
Processes acquire locks in same order

锁释放的顺序无关紧要 Order in which locks released immaterial

线程0 Thread 0

$s_0=s_1=1$



演示 Demonstration

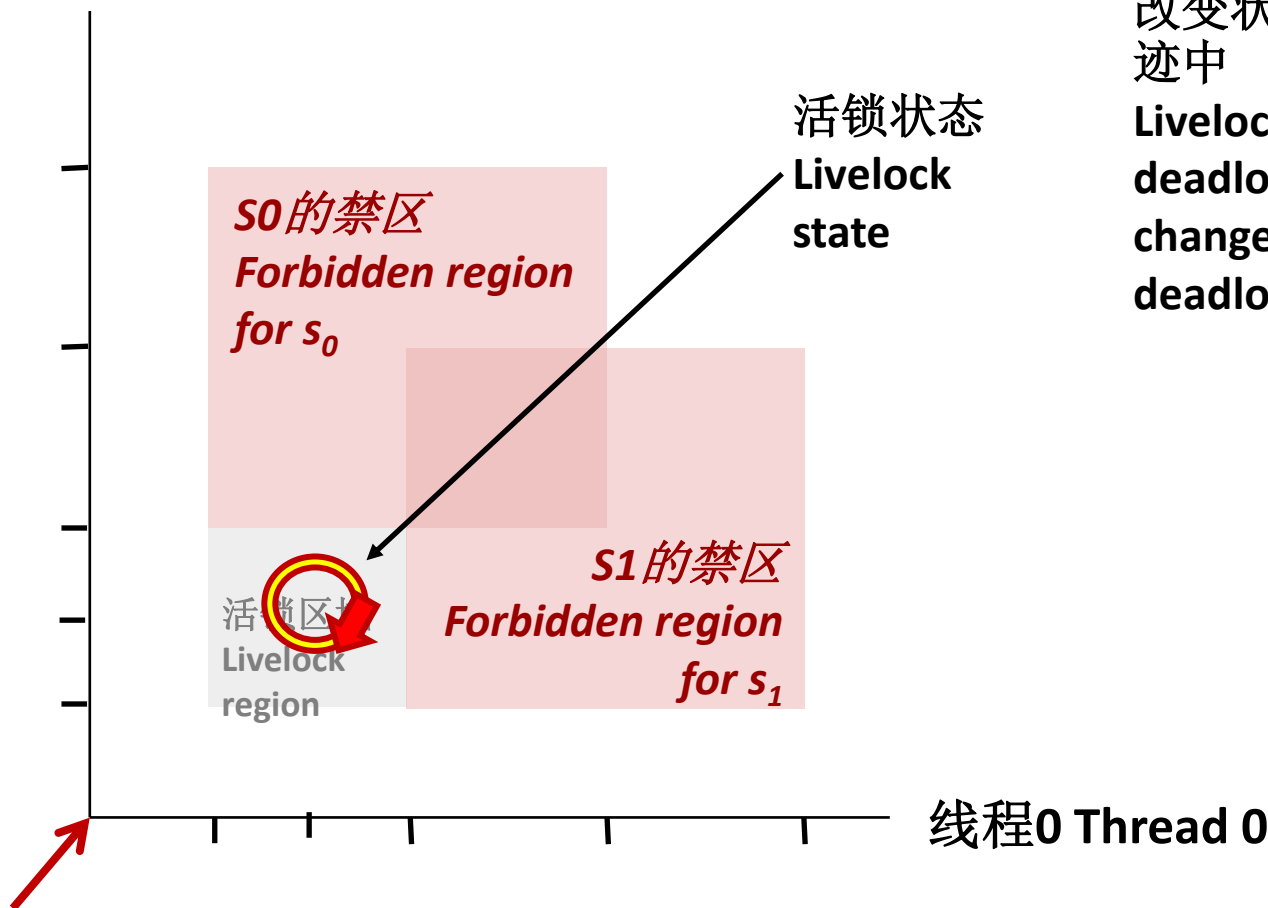
- 参见程序 `deadlock.c` See program `deadlock.c`
- 100个线程，每个线程获得同样的两个锁 100 threads, each acquiring same two locks
- 风险模式 Risky mode
 - 偶数线程请求锁的顺序与奇数线程相反 Even numbered threads request locks in opposite order of odd-numbered ones
- 安全模式 Safe mode
 - 所有线程以相同的顺序获取锁 All threads acquire locks in same order

在进度图中显示活锁

Livelock Visualized in Progress Graph



线程1 Thread 1



活锁类似于死锁，只是线程改变状态，但仍处于死锁轨迹中

Livelock is similar to a deadlock, except the threads change state, but remain in a deadlock trajectory.

死锁、活锁、饿死

Deadlock, Livelock, Starvation



■ 死锁 **Deadlock**

- 一个或多个线程正在等待一个永远不会为真的条件 One or more threads is waiting on a condition that will never be true

■ 活锁 **Livelock**

- 一个或多个线程正在更改状态，但永远不会离开死锁/活锁轨迹
One or more threads is changing state, but will never leave a deadlock / livelock trajectory

■ 饿死 **Starvation**

- 一个或多个线程暂时无法取得进展 One or more threads is temporarily unable to make progress



议题 Today

- 回顾：信号量、互斥和生产者-消费者 Review: Semaphores, mutexes, producer-consumer
- 使用信号量调度共享资源 Using semaphores to schedule shared resources
 - 读者-写者问题 Readers-writers problem
- **其它并发问题 Other concurrency issues**
 - 竞争 Races
 - 死锁 Deadlocks
 - **线程安全 Thread safety**
 - 线程和信号处理之间交互 Interactions between threads and signal handling

关键概念：线程安全

Crucial concept: Thread Safety



- 从线程调用的函数必须是**线程安全**的 Functions called from a thread must be **thread-safe**
- 定义：函数是线程安全的，只要它在从多个线程同时调用时总是产生正确的结果 **Def: A function is *thread-safe* iff it will always produce correct results when called simultaneously from multiple threads.**
- 线程不安全函数的分类： **Classes of thread-unsafe functions:**
 - 类1：不保护共享变量的函数 Class 1: Functions that do not protect shared variables
 - 类2：跨多个调用保持状态的函数 Class 2: Functions that keep state across multiple invocations
 - 类3：返回指向静态变量的指针的函数 Class 3: Functions that return a pointer to a static variable
 - 类4：调用线程不安全函数的函数 Class 4: Functions that call thread-unsafe functions

线程不安全函数（类1）

Thread-Unsafe Functions (Class 1)



- 未能保护共享变量 **Failing to protect shared variables**
 - 修复：使用P和V信号量操作（或互斥锁） Fix: Use *P* and *V* semaphore operations (or mutex)
 - 示例：goodcnt.c Example: **goodcnt.c**
 - 问题：同步操作会降低代码速度 Issue: Synchronization operations will slow down code

线程不安全函数（类2）



Thread-Unsafe Functions (Class 2)

- 跨多个函数调用依赖持久状态 Relying on persistent state across multiple function invocations
 - 示例：依赖于静态状态的随机数生成器 Example: Random number generator that relies on static state

```
static unsigned int next = 1;

/* rand: return pseudo-random integer on 0..32767 */
int rand(void)
{
    next = next*1103515245 + 12345;
    return (unsigned int)(next/65536) % 32768;
}

/* srand: set seed for rand() */
void srand(unsigned int seed)
{
    next = seed;
}
```



线程安全随机数生成器

Thread-Safe Random Number Generator

- 传递状态作为参数的一部分 **Pass state as part of argument**
 - 从而消除静态状态 and, thereby, eliminate static state

```
/* rand_r - return pseudo-random integer on 0..32767 */  
  
int rand_r(int *nextp)  
{  
    *nextp = *nextp*1103515245 + 12345;  
    return (unsigned int)(*nextp/65536) % 32768;  
}
```

- 结论：使用rand_r的程序员必须保持种子 **Consequence:**
programmer using rand_r must maintain seed

线程不安全函数（类3）

Thread-Unsafe Functions (Class 3)



- 返回指向静态变量的指针
Returning a pointer to a static variable
- 修复：重写函数，以便调用方传递变量地址以存储结果 **Fix: Rewrite function so caller passes address of variable to store result**
 - 需要更改调用者和被调用者
Requires changes in caller and callee
- 修复2：用互斥锁包装函数 **Fix 2: Wrap function with mutex**
 - 调用方仍需更改 Caller still has to be changed
 - 可以保留旧函数 Can preserve old function
 - 函数可能成为瓶颈 Function may become a bottleneck

```
/* Convert integer to string */
char *itoa(int x)
{
    static char buf[11];
    snprintf(buf, sizeof buf,
              "%d", x);
    return buf;
}
```

```
void fix_itoa(int x, char *dst,
              size_t dstsz)
{
    snprintf(dst, dstsz, "%d", x);
}
```

```
void wrap_itoa(int x, char *dst,
               size_t dstsz)
{
    static sem_t mutex;
    P(mutex);
    strncpy(dst, itoa(x), dstsz);
    V(mutex);
}
```

线程不安全函数（类4）



Thread-Unsafe Functions (Class 4)

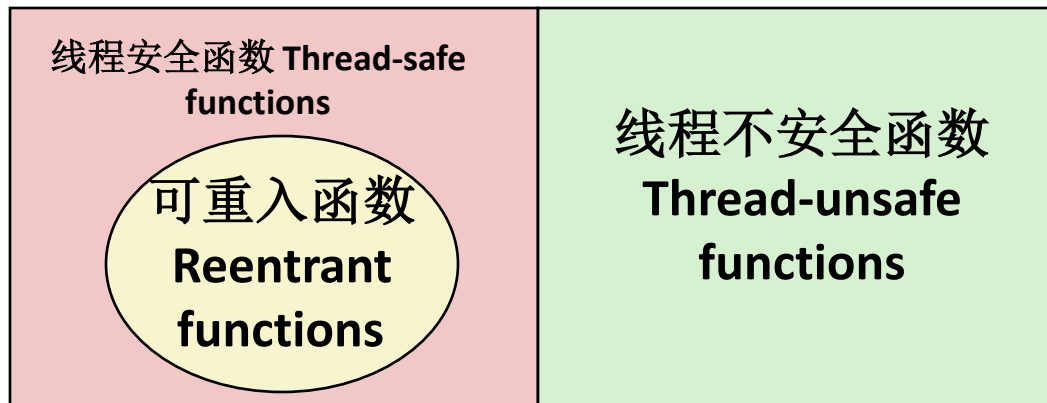
- 调用线程不安全函数 **Calling thread-unsafe functions**
 - 调用一个线程不安全函数会使调用它的整个函数不安全 Calling one thread-unsafe function makes the entire function that calls it thread-unsafe
 - 修复：修改函数，使其仅调用线程安全函数 Fix: Modify the function so it calls only thread-safe functions 😊



可重入函数 Reentrant Functions

- 定义：当且仅当函数被多个线程调用时不访问共享变量，则该函数是 **可重入** 的 Def: A function is **reentrant** iff it accesses no shared variables when called by multiple threads.
 - 线程安全函数的重要子集 Important subset of thread-safe functions
 - 不需要同步操作 Require no synchronization operations
 - 使类2函数线程安全的唯一方法是使其可重入（例如rand_r）
Only way to make a Class 2 function thread-safe is to make it reentrant (e.g., rand_r)

全部函数 All functions





线程安全的库函数

Thread-Safe Library Functions

- 标准C语言库（K&R教材后面）中的大多数函数都是线程安全的
Most functions in the Standard C Library (at the back of your K&R text) are thread-safe
 - 示例：malloc、free、printf、scanf Examples: **malloc**, **free**, **printf**, **scanf**
 - 例外：strtok、rand、ctime Exceptions: **strtok**, **rand**, **ctime**
- POSIX添加了更多的异常，但也添加不安全函数的可重入版本
POSIX adds more exceptions, but also reentrant versions of unsafe functions


线程不安全函数 Thread-unsafe function	Class	可重入版 Reentrant version
asctime	3	strptime
ctime	3	strptime
gethostbyaddr	3	getnameinfo
gethostbyname	3	getaddrinfo
inet_ntoa	3	getnameinfo
localtime	3	localtime_r
rand	2	rand_r



议题 Today

- 回顾：信号量、互斥和生产者-消费者 Review: Semaphores, mutexes, producer-consumer
- 使用信号量调度共享资源 Using semaphores to schedule shared resources
 - 读者-写者问题 Readers-writers problem
- **其它并发问题 Other concurrency issues**
 - 竞争 Races
 - 死锁 Deadlocks
 - 线程安全 Thread safety
 - **线程和信号处理之间交互** Interactions between threads and signal handling

信号处理回顾 Signal Handling Review




■ 动作 Action

- 信号可以发生在程序执行的任何点 Signal can occur at any point in program execution
 - 除非信号被阻塞 Unless signal is blocked
- 信号处理程序在同一个线程内运行 Signal handler runs within same thread
- 必须运行到完成，然后返回到正常的程序执行 Must run to completion and then return to regular program execution

线程/信号交互


Threads / Signals Interactions



- 很多库函数有内部锁 Many library functions have internal locking
 - 为了保护隐藏状态（避免第一类线程不安全） To protect hidden state (avoid being class 1 thread-unsafe)
 - malloc
 - 释放列表 Free lists
 - fputs, fprintf, snprintf
 - 以便从多个线程的输出不会交错 So that outputs from multiple threads don't interleave
 - 内部使用malloc Internal use of malloc
- 不使用这些库函数的处理程序没有问题 OK for handler that doesn't use these library functions

有问题的线程/信号交互

Bad Thread / Signal Interactions



■ 如果以下情况会怎样: What if:

- 当库函数保持加锁时接收信号 Signal received while library function holds lock
- 处理程序调用同样 (或相关) 库函数 Handler calls same (or related) library function

■ 死锁! Deadlock!

- 信号处理程序不能继续直到获得锁 Signal handler cannot proceed until it gets lock
- 主程序不能继续直到处理程序完成 Main program cannot proceed until handler completes

■ 关键点 Key Point

- 线程采用对称并发 Threads employ symmetric concurrency
- 信号处理是异步的 Signal handling is asymmetric

处理线程/信号交互

Handling Thread/Signal Interactions



- **临界区周围阻塞信号 Block signals around critical sections**
 - pthread_sigmask函数类似sigprocmask, 但是仅影响正调用的线程
pthread_sigmask – like sigprocmask, but only affects calling thread
- **专用于信号处理的线程 Dedicate a thread to signal handling**
 - 循环调用sigsuspend()或sigwaitinfo() Loop calling sigsuspend() or sigwaitinfo()
 - 所有其他线程阻塞所有信号 All other threads block all signals
 - 信号处理线程可以使用异步信号不安全函数 Signal handling thread can use async-signal-unsafe functions
 - 因为我们知道信号只能在sigsuspend()期间传递 Because we know signals will only be delivered during sigsuspend()

线程小结 Threads Summary



- 线程为编写并发程序提供了另一种机制 **Threads provide another mechanism for writing concurrent programs**
- 线程越来越受欢迎 **Threads are growing in popularity**
 - 比进程开销小 Somewhat cheaper than processes
 - 易于在线程之间共享数据 Easy to share data between threads
- 然而，共享的便捷性有代价： **However, the ease of sharing has a cost:**
 - 易于引入细微的同步错误 Easy to introduce subtle synchronization errors
 - 小心对待线程！ Tread carefully with threads!
- 有关详细信息： **For more info:**
 - “Posix线程编程” D. Butenhof, “Programming with Posix Threads”, Addison-Wesley, 1997