

Database Principles and Design

(数据库原理与设计)

Xiaolin Zhao

Chapter I:

The Worlds of Database Systems

Outline

- Why do we need a DBMS (Database Management System)?
- What can a DBMS do for an application?
- Why study database systems?
- Data Models: Overview of a Relational Model
- Levels of Abstraction in a DBMS
- Sample Queries in DBMS
- Transaction Management Overview
- Structure of a DBMS

Why DBMS?

- Suppose that you want to build an university database. It must store the following information:
 - Entities: Students, Professors, Classes, Classrooms
 - Relationships: Who teaches what? Who teaches where? Who teaches whom?

What can DBMS do for applications?

- Store huge amount of data (e.g., TB+) over a long period of time
- Allow apps to query and update data
 - Query: what is Mary's grade in the "Operating System" course?
 - Update: enroll Mary in the "Database" course
- Protect from unauthorized access.
 - Students cannot change their course grades.
- Protect from system crashes
 - When some system components fail (hard drive, network, etc.), database can be restored to a good state.

More on what can DBMS do for applications?

- Protect from incorrect inputs
 - Mary has registered for 100 courses
- Support concurrent access from multiple users
 - 1000 students using the registration system at the same time
- Allow administrators to easily change data schema
 - At a later time, add TA info to courses.
- Efficient database operations
 - Search for students with 5 highest GPAs

Alternative to Using a DBMS

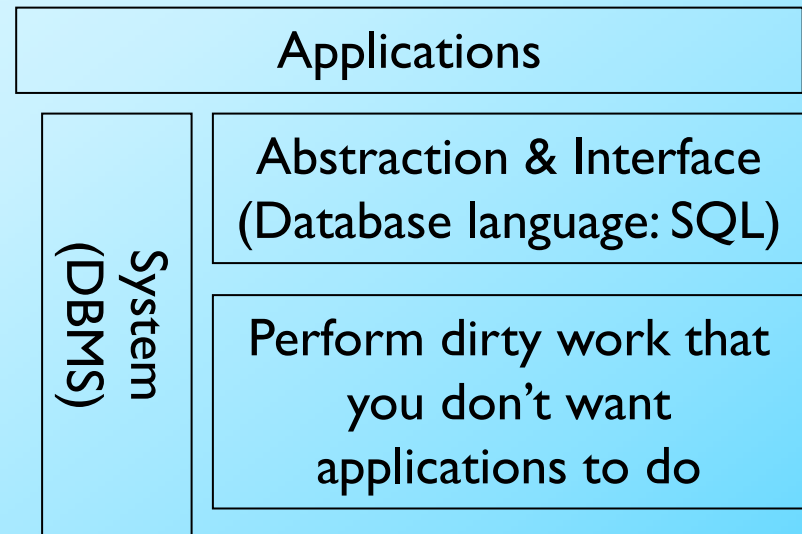
- Store data as files in operating systems.
- Applications have to deal with the following issues:
 - 32-bit addressing (4GB) is insufficient to address 100GB+ data file
 - Write special code to support different queries
 - Write special code to protect data from concurrent access
 - Write special code to protect against system crashes
 - Optimize applications for efficient access and query
 - May often rewrite applications
- Easier to buy a DBMS to handle these issues

Database Management System (DBMS)

- DBMS is software to store and manage data, so applications don't have to worry about them.
- What can a DBMS do for applications?
 - Can you think of them?

What can a DBMS do for applications?

- Define data: Data Definition Language (DDL)
- Access and operate on data: Data Manipulation Language (DML)
 - Query language
- Storage management
- Transaction Management
 - Concurrency control
 - Crash recovery
- Provide good security, efficiency, and scalability



Why Study Database Systems?

- They are everywhere.
 - Online stores, real stores
 - Banks, credit card companies
 - Passport control
 - Police (criminal records)
 - Airlines and hotels (reservations)
- DBMS vendors & products
 - Oracle, Microsoft (Access and SQL server), IBM (DB2), Sybase, ...

Data Models

- A **data model** is a collection of concepts for describing data.
 - Entity-relation (ER) model
 - Relational model (main focus of this course)
- A **schema** is a description of data.
- The **relational model** is the most widely used data model.
 - A **relation** is basically a table with rows and columns of **records**.
 - Every relation has a **schema**, which describes the columns, or fields.

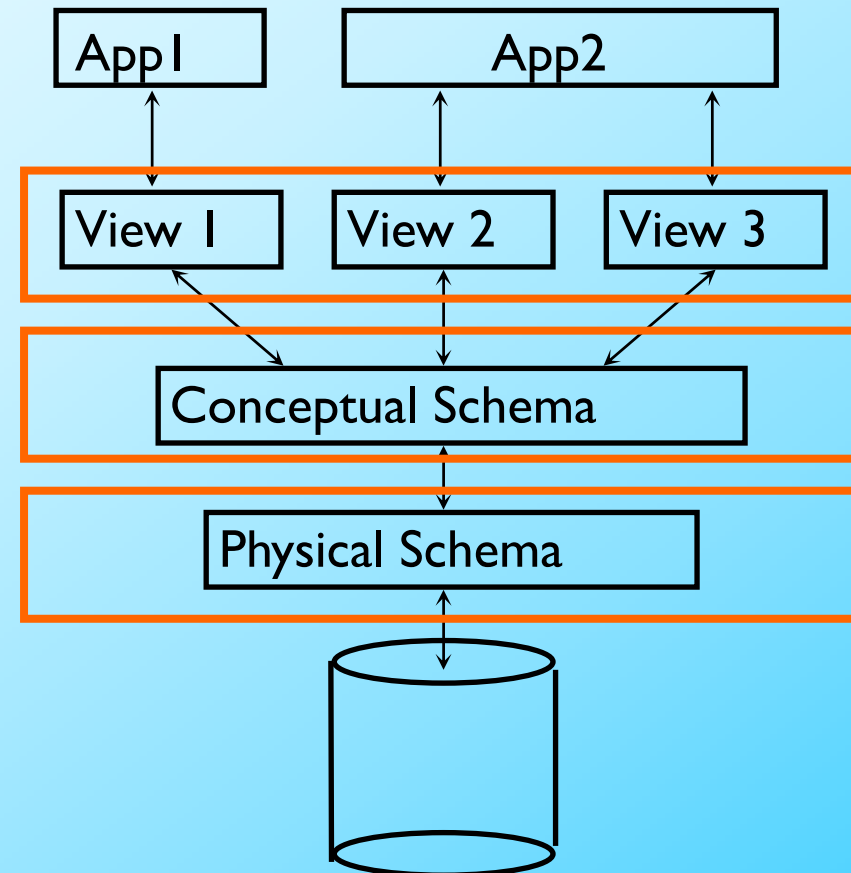
Relational Model

- The entire table shows an instance of the Students relation.
- The Students schema is the column heads
 - Students(Sid: String, Name: String, Login: String, age: Integer,...)

sid	name	email	age	gpa
53666	Jones	Jones@cs	18	3.4
53688	Smith	Smith@ee	18	3.2
53650	Joe	Joe@cs	19	2.5

Levels of Abstractions in DBMS

- Many **views**, one **conceptual schema** and one **physical schema**.
 - Conceptual schema defines logical structure
 - Relation tables
 - Physical schema describes the file and indexing used
 - Sorted file with B+ tree index
 - Views describe how applications (users) see the data
 - Relation tables but not store explicitly

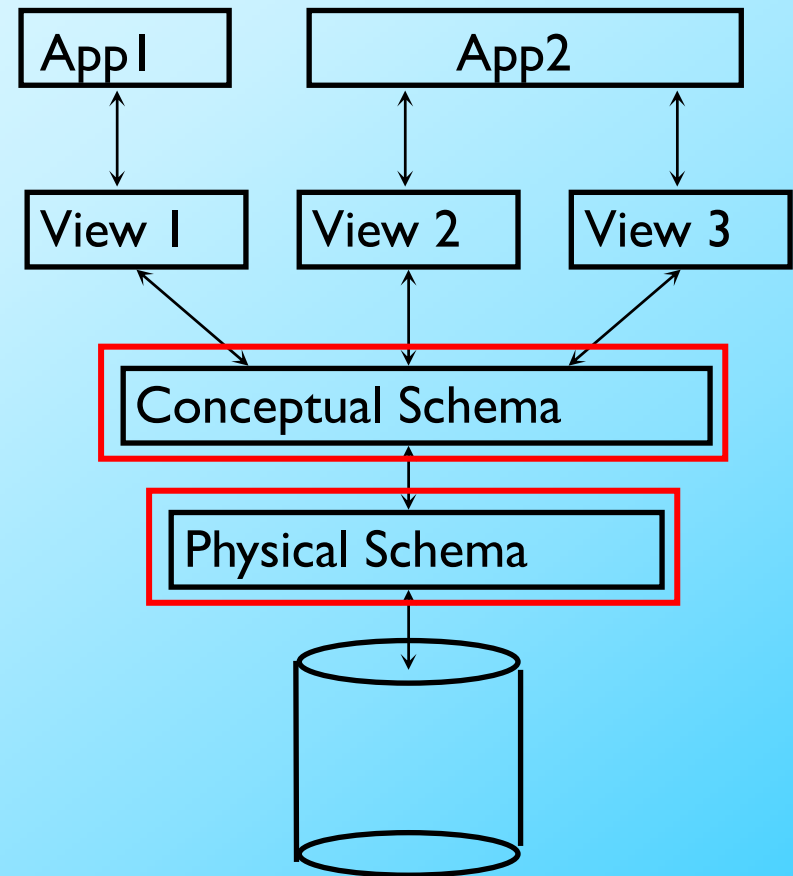


Example: University Database

- Conceptual schema:
 - Students (sid: string, name: string, login: string, age: integer, gpa:real)
 - Courses (cid: string, cname:string, credits:integer)
 - Enrolled (sid:string, cid:string, grade:string)
- Physical schema:
 - Relations stored as unordered files.
 - Index on first column of Students.
- View (External Schema):
 - Course_info(cid:string, enrollment:integer)
 - Why?

Data Independence

- Three levels of abstraction provides **data independence**.
 - Changes in one layer only affect one upper layer.
 - E.g., applications are not affected by changes in conceptual & physical schema.



Queries in DBMS

- Sample queries on university database:
 - What is the name of the student with student ID 123456?
- The key benefits of using a relational database are
 - Easy to specify queries using a **query language**: Structured Query Language (SQL)

```
SELECT S.name  
FROM Students S  
WHERE S.sid = 123456
```
 - Efficient query processor to get answer

Transaction Management

- A transaction is an execution of a user program in a DBMS.
- Transaction management deals with two things:
 - Concurrent execution of transactions
 - Incomplete transactions and system crashes

Concurrency Control

- Example: two travel agents (A, B) are trying to book one remaining airline seat (two transactions), only one transaction can succeed in booking.

`// num_seats is 1`

Transactions A and B: if num_seats > 0, book the seat & num_seat--;

`// overbook!`

- How to solve this?

Concurrency Control (Solution)

// num_seats is 1

Transactions A and B: if num_seats > 0, book the seat & num_seat--;
// overbook!

- **Solution: use locking protocol**

Transaction A: get exclusive lock on **num_seats**

Transaction B: wait until A releases lock on num_seats

Transaction A: if num_seats > 0, book & num_seat--;

// book the seat, num_seat is

set to 0

Transaction A: release exclusive lock on num_seats

Transaction B: num_seats = 0, no booking; // does not book the seat

Crash Recovery

- Example: a bank transaction transfers \$100 from account A to account B

$A = A - \$100$

<system crashes> // good for the bank!

$B = B + \$100$

- How to solve this?

Crash Recovery (Solution)

A = A - \$100

<system crashes> // good for the bank!

B = B + \$100

- Solution: use logging, meaning that all write operations are recorded in a log on a stable storage.

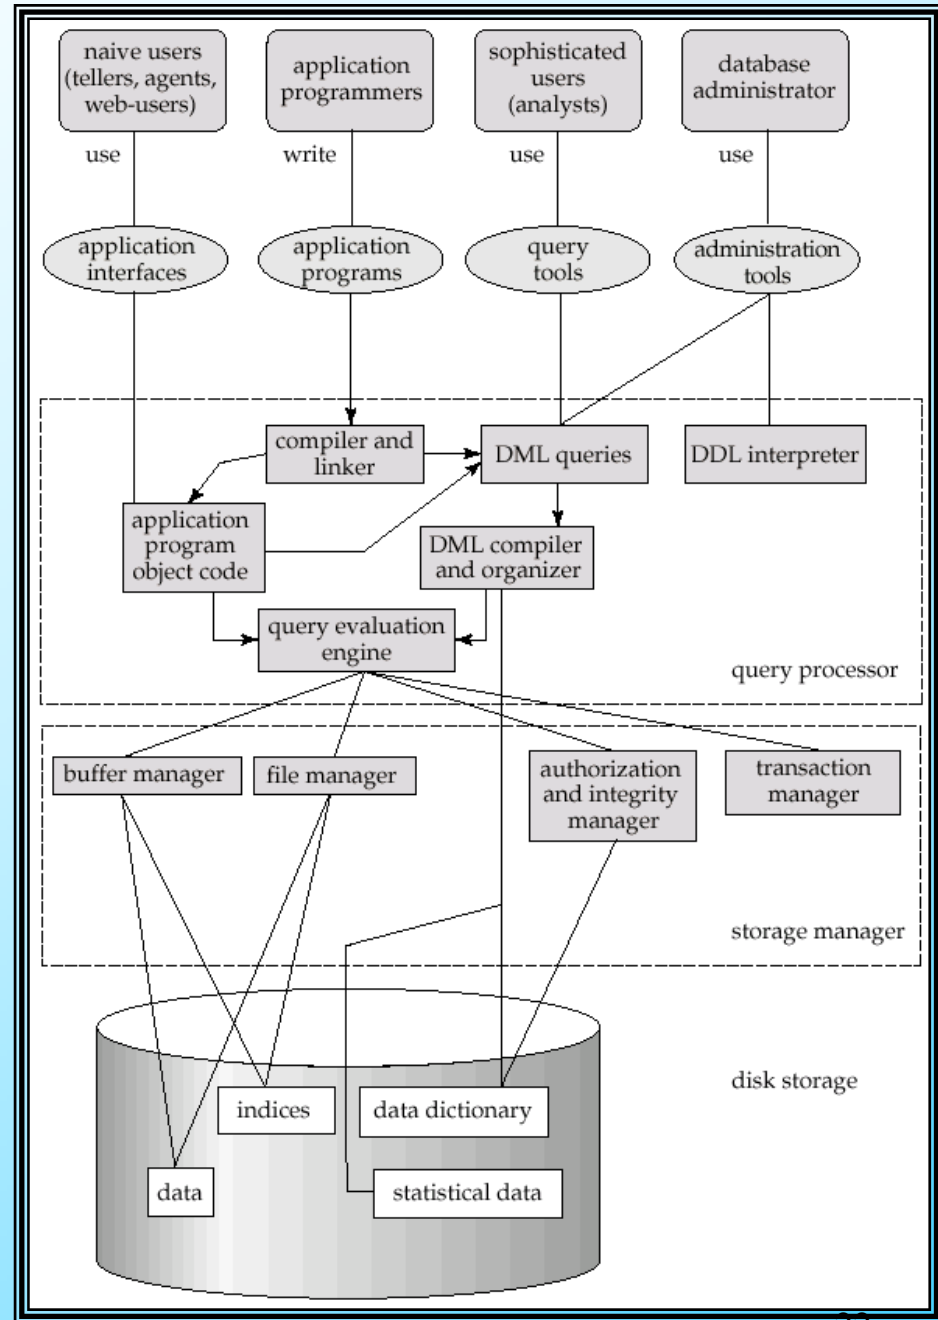
A = A - \$100 // recorded A value (checkpoint) in a log

<system crashes>

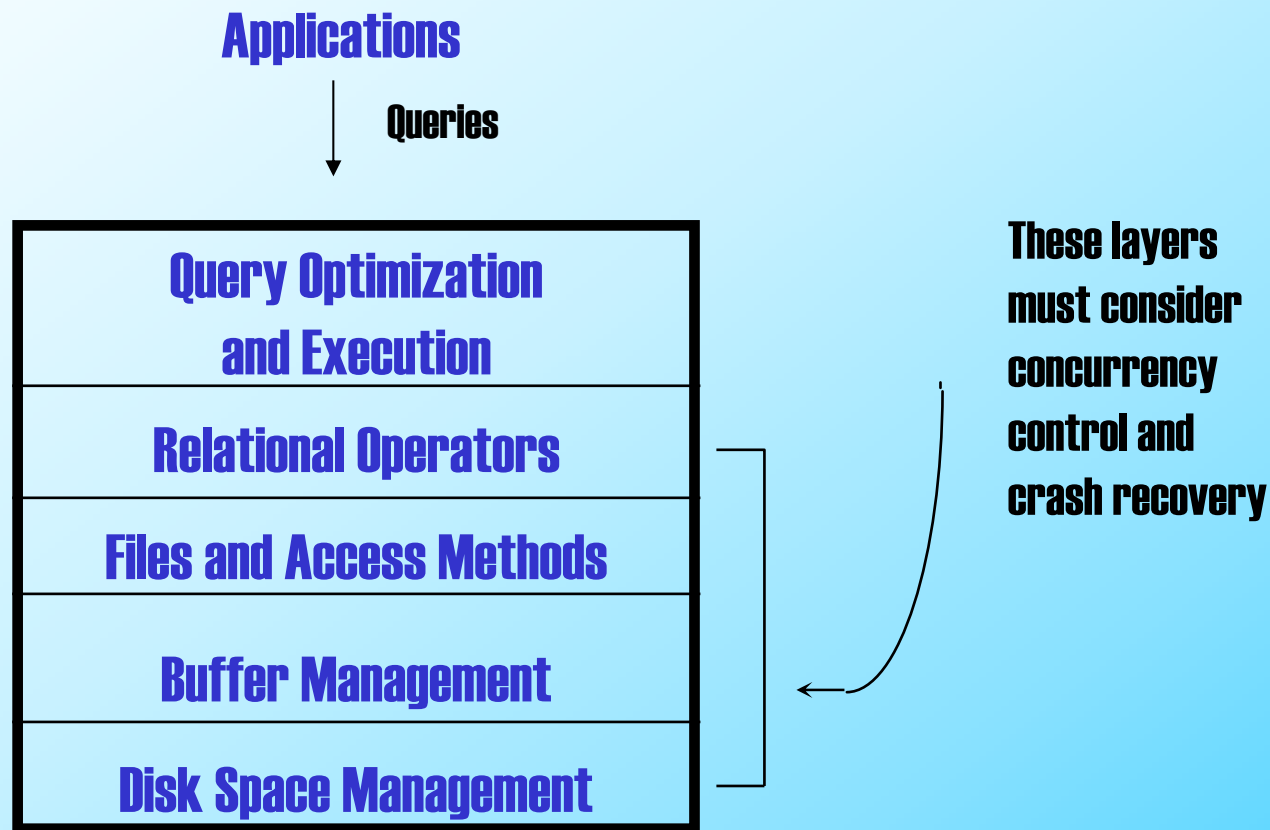
// start recovery: read the log from disk

//analyze, undo, & redo

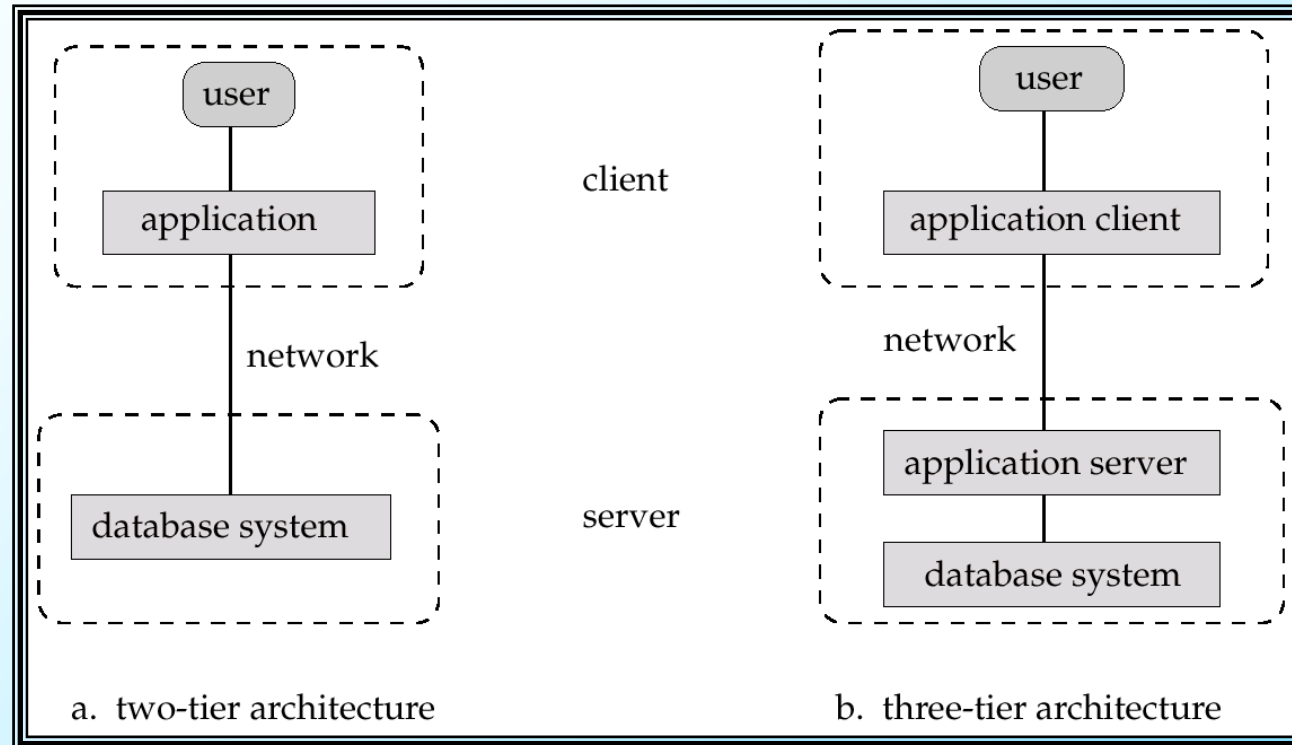
Overview of DBMS



Layered Architecture



Application Architectures



- **Two-tier architecture:** E.g. client programs using ODBC/JDBC to communicate with a database
- **Three-tier architecture:** E.g. web-based applications, and applications built using “middleware”

Homework

- Read Chapters 1
- Read Chapter 2 for next lecture