

Real SQL Programming

Persistent Stored Modules (PSM)

PL/SQL

Embedded SQL

SQL in Real Programs

- ◆ We have seen only how SQL is used at the generic query interface --- an environment where we sit at a terminal and ask queries of a database.
- ◆ Reality is almost always different: conventional programs interacting with SQL.

Options

1. Code in a specialized language is stored in the database itself (e.g., PSM, PL/SQL).
2. SQL statements are embedded in a *host language* (e.g., C).
3. Connection tools are used to allow a conventional language to access a database (e.g., CLI, JDBC, PHP/DB).

Control-of-Flow Language Elements

Example

◆ Statement Level

- ◆ BEGIN ... END block
- ◆ IF ... ELSE block
- ◆ WHILE constructs

◆ Row Level

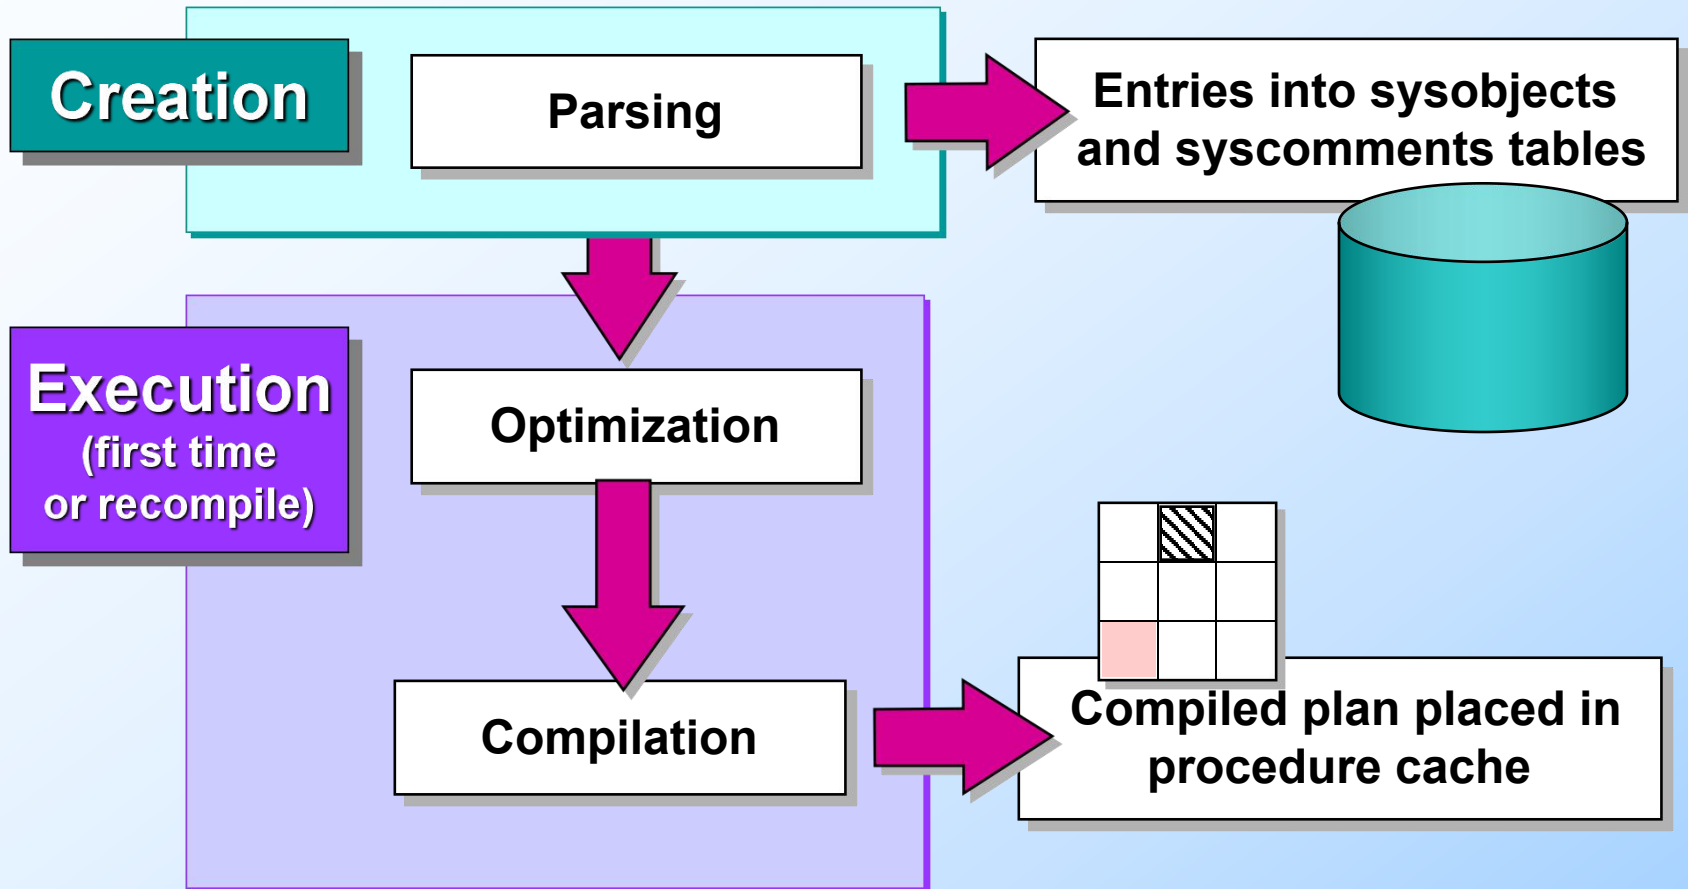
- ◆ CASE expression

```
DECLARE @n tinyint
SET @n = 5
IF (@n BETWEEN 4 and 6)
BEGIN
    WHILE (@n > 0)
    BEGIN
        SELECT  @n AS 'Number'
                ,CASE
                    WHEN (@n % 2) = 1
                    THEN 'EVEN'
                    ELSE 'ODD'
                END AS 'Type'
        SET @n = @n - 1
    END
END
ELSE
    PRINT 'NO ANALYSIS'
```

Stored Procedures

- ◆ PSM, or “*persistent stored modules*,” allows us to store procedures as database schema elements.
- ◆ PSM = a mixture of conventional statements (if, while, etc.) and SQL.
- ◆ Lets us do things we cannot do in SQL alone.

Initial Processing of Stored Procedures



Advantages of Stored Procedures

- ◆ Share Application Logic
- ◆ Shield Database Schema Details
- ◆ Provide Security Mechanisms
- ◆ Improve Performance
- ◆ Reduce Network Traffic

Basic PSM Form

```
CREATE PROCEDURE <name> (  
    <parameter list> )  
    <optional local declarations>  
    <body>;
```

◆ Function alternative:

```
CREATE FUNCTION <name> (  
    <parameter list> ) RETURNS <type>
```


Parameters in PSM

- ◆ Unlike the usual name-type pairs in languages like C, PSM uses mode-name-type triples, where the *mode* can be:
 - ◆ IN = procedure uses value, does not change value.
 - ◆ OUT = procedure changes, does not use.
 - ◆ INOUT = both.

Example: Stored Procedure

- ◆ Let's write a procedure that takes two arguments b and p , and adds a tuple to `Sells(bar, beer, price)` that has `bar = 'Joe's Bar'`, `beer = b` , and `price = p` .
 - ◆ Used by Joe to add to his menu more easily.

The Procedure

```
CREATE PROCEDURE JoeMenu (
```

```
IN b    CHAR(20),  
IN p    REAL
```

Parameters are both
read-only, not changed

```
)
```

```
INSERT INTO Sells  
VALUES('Joe"s Bar', b, p);
```

The body ---
a single insertion

Invoking Procedures

- ◆ Use SQL/PSM statement `CALL`, with the name of the desired procedure and arguments.

- ◆ **Example:**

```
CALL JoeMenu('Moosedrool', 5.00);
```

- ◆ Functions used in SQL expressions wherever a value of their return type is appropriate.

Kinds of PSM statements – (1)

- ◆ RETURN <expression> sets the return value of a function.
 - ◆ Unlike C, etc., RETURN *does not* terminate function execution.
- ◆ DECLARE <name> <type> used to declare local variables.
- ◆ BEGIN . . . END for groups of statements.
 - ◆ Separate statements by semicolons.

Kinds of PSM Statements – (2)

◆ Assignment statements:

SET <variable> = <expression>;

◆ Example: SET b = 'Bud' ;

◆ Statement labels: give a statement a label by prefixing a name and a colon.

IF Statements

- ◆ Simplest form:

```
IF <condition> THEN  
    <statements(s)>  
END IF;
```

- ◆ Add ELSE <statement(s)> if desired, as
IF . . . THEN . . . ELSE . . . END IF;

- ◆ Add additional cases by ELSEIF
<statements(s)>: IF ... THEN ... ELSEIF ...
THEN ... ELSEIF ... THEN ... ELSE ... END IF;

Example: IF

- ◆ Let's rate bars by how many customers they have, based on `Frequents(drinker,bar)`.
 - ◆ < 100 customers: 'unpopular'.
 - ◆ 100-199 customers: 'average'.
 - ◆ ≥ 200 customers: 'popular'.
- ◆ Function `Rate(b)` rates bar b.

Example: IF (continued)

```
CREATE FUNCTION Rate (IN b CHAR(20) )
```

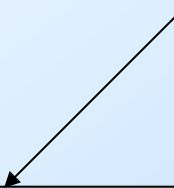
```
  RETURNS CHAR(10)
```

```
  DECLARE cust INTEGER;
```

```
  BEGIN
```

```
    SET cust = (SELECT COUNT(*) FROM Frequents  
                WHERE bar = b);
```

Number of
customers of
bar b



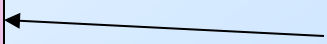
```
    IF cust < 100 THEN RETURN 'unpopular'  
    ELSEIF cust < 200 THEN RETURN 'average'  
    ELSE RETURN 'popular'  
    END IF;
```

Nested
IF statement



```
  END;
```

Return occurs here, not at
one of the RETURN statements



Loops

- ◆ Basic form:

```
<loop name>: LOOP <statements>  
                END LOOP;
```

- ◆ Exit from a loop by:

```
    LEAVE <loop name>
```

Example: Exiting a Loop

```
loop1: LOOP
```

```
  . . .
```

```
  LEAVE loop1; ← If this statement is executed . . .
```

```
  . . .
```

```
END LOOP;
```

```
← Control winds up here
```

Other Loop Forms

- ◆ WHILE <condition>
 DO <statements>
 END WHILE;
- ◆ REPEAT <statements>
 UNTIL <condition>
 END REPEAT;

Executing Stored Procedures

◆ Executing a Stored Procedure by Itself

```
EXEC OverdueOrders
```

◆ Executing a Stored Procedure Within an INSERT Statement

```
INSERT INTO Customers  
EXEC EmployeeCustomer
```

Using Input Parameters

- ◆ Validate All Incoming Parameter Values First
- ◆ Provide Appropriate Default Values and Include Null Checks

```
CREATE PROCEDURE dbo.[Year to Year Sales]
    @BeginningDate DateTime, @EndingDate DateTime
AS
IF @BeginningDate IS NULL OR @EndingDate IS NULL
BEGIN
    RAISERROR('NULL values are not allowed', 14, 1)
    RETURN
END
SELECT O.ShippedDate,
       O.OrderID,
       OS.Subtotal,
       DATENAME(yy, ShippedDate) AS Year
FROM ORDERS O INNER JOIN [Order Subtotals] OS
    ON O.OrderID = OS.OrderID
WHERE O.ShippedDate BETWEEN @BeginningDate AND @EndingDate
```

Executing Stored Procedures Using Input Parameters

◆ Passing Values by Parameter Name

```
EXEC AddCustomer  
    @CustomerID = 'ALFKI',  
    @ContactName = 'Maria Anders',  
    @CompanyName = 'Alfreds Futterkiste',  
    @ContactTitle = 'Sales Representative',  
    @Address = 'Obere Str. 57',  
    @City = 'Berlin',  
    @PostalCode = '12209',  
    @Country = 'Germany',  
    @Phone = '030-0074321'
```

◆ Passing Values by Position

```
EXEC AddCustomer 'ALFKI2', 'Alfreds  
Futterkiste', 'Maria Anders', 'Sales  
Representative', 'Obere Str. 57', 'Berlin',  
NULL, '12209', 'Germany', '030-0074321'
```

Returning Values Using Output Parameters

Creating Stored Procedure

```
CREATE PROCEDURE dbo.MathTutor  
    @m1 smallint,  
    @m2 smallint,  
    @result smallint OUTPUT  
AS  
    SET @result = @m1* @m2
```

Executing Stored Procedure

```
DECLARE @answer smallint  
EXECUTE MathTutor 5,6, @answer OUTPUT  
SELECT 'The result is: ', @answer
```

Results of Stored Procedure

```
The result is: 30
```


Queries

- ◆ General SELECT-FROM-WHERE queries are *not* permitted in PSM.
- ◆ There are three ways to get the effect of a query:
 1. Queries producing one value can be the expression in an assignment.
 2. Single-row SELECT . . . INTO.
 3. Cursors.

Example: Assignment/Query

- ◆ Using local variable p and `Sells(bar, beer, price)`, we can get the price Joe charges for Bud by:

```
SET p = (SELECT price FROM Sells
        WHERE bar = 'Joe''s Bar' AND
               beer = 'Bud' );
```

SELECT . . . INTO

- ◆ Another way to get the value of a query that returns one tuple is by placing **INTO** **<variable>** after the SELECT clause.

- ◆ **Example:**

```
SELECT price INTO p FROM Sells
WHERE bar = 'Joe''s Bar' AND
      beer = 'Bud';
```

Cursors

◆ A *cursor* is essentially a tuple-variable that ranges over all tuples in the result of some query.

◆ Declare a cursor *c* by:

```
DECLARE c CURSOR FOR <query>;
```

Opening and Closing Cursors

- ◆ To use cursor c , we must issue the command:

OPEN c ;

- ◆ The query of c is evaluated, and c is set to point to the first tuple of the result.

- ◆ When finished with c , issue command:

CLOSE c ;

Fetching Tuples From a Cursor

- ◆ To get the next tuple from cursor c , issue command:

`FETCH FROM c INTO x_1, x_2, \dots, x_n ;`

- ◆ The x 's are a list of variables, one for each component of the tuples referred to by c .
- ◆ c is moved automatically to the next tuple.

Breaking Cursor Loops – (1)

- ◆ The usual way to use a cursor is to create a loop with a FETCH statement, and do something with each tuple fetched.
- ◆ A tricky point is how we get out of the loop when the cursor has no more tuples to deliver.

Breaking Cursor Loops – (2)

- ◆ Each SQL operation returns a *status*, which is a 5-digit character string.
 - ◆ For example, 00000 = "Everything OK," and 02000 = "Failed to find a tuple."
- ◆ In PSM, we can get the value of the status in a variable called SQLSTATE.

Breaking Cursor Loops – (3)

◆ We may declare a *condition*, which is a boolean variable that is true if and only if SQLSTATE has a particular value.

◆ **Example:** We can declare condition NotFound to represent 02000 by:

```
DECLARE NotFound CONDITION FOR  
SQLSTATE '02000';
```

Breaking Cursor Loops – (4)

◆ The structure of a cursor loop is thus:

```
cursorLoop: LOOP
```

```
...
```

```
FETCH c INTO ... ;
```

```
IF NotFound THEN LEAVE
```

```
cursorLoop;
```

```
END IF;
```

```
...
```

```
END LOOP;
```

Example: Cursor

- ◆ Let's write a procedure that examines `Sells(bar, beer, price)`, and raises by \$1 the price of all beers at Joe's Bar that are under \$3.
 - ◆ Yes, we could write this as a simple UPDATE, but the details are instructive anyway.

The Needed Declarations

```
CREATE PROCEDURE JoeGouge( )
```

```
    DECLARE theBeer CHAR(20);
```

```
    DECLARE thePrice REAL;
```

Used to hold
beer-price pairs
when fetching
through cursor c

```
    DECLARE NotFound CONDITION FOR  
        SQLSTATE '02000';
```

```
    DECLARE c CURSOR FOR
```

```
        (SELECT beer, price FROM Sells  
         WHERE bar = 'Joe's Bar');
```

Returns Joe's menu

The Procedure Body


BEGIN

OPEN c;

menuLoop: LOOP

FETCH c INTO theBeer, thePrice;

Check if the recent
FETCH failed to
get a tuple



IF NotFound THEN LEAVE menuLoop END IF;

IF thePrice < 3.00 THEN

UPDATE Sells SET price = thePrice + 1.00

WHERE bar = 'Joe's Bar' AND beer = theBeer;

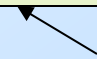
END IF;

END LOOP;

CLOSE c;

END;

If Joe charges less than \$3 for
the beer, raise its price at
Joe's Bar by \$1.



游标的运行

1	EXEC SQL BEGIN //定义主变量	2	:PNAME	:PXH
1	EXEC SQL DECLARE tnames_cursor CURSOR	3	张三	003
	FOR SELECT name,xh FROM xjb		李四	004
	WHERE xb='男'		王五	005
2	ORDER BY xh;		孙六	006
3	EXEC SQL OPEN tnames_cursor;		陈七	007
3	EXEC SQL FETCH tnames_cursor INTO :PNAME,:PXH;		曹八	008
	WHILE (:PXH!=NULL)		.	
	{		.	
4	printf("%s,%s\n",PNAME,PXH);		.	
	EXEC SQL FETCH tnames_cursor INTO :PNAME,:PXH			
	}			
5	EXEC SQL CLOSE tnames_cursor;			
	EXEC SQL DEALLOCATE tnames_cursor;			

4 Results

张三	003
李四	004
王五	005
孙六	006
陈七	007
曹八	008
...	

PL/SQL

- ◆ Oracle uses a variant of SQL/PSM which it calls PL/SQL.
- ◆ PL/SQL not only allows you to create and store procedures or functions, but it can be run from the *generic query interface* (sqlplus), like any SQL statement.
- ◆ Triggers are a part of PL/SQL.

Trigger Differences

- ◆ Compared with SQL standard triggers, Oracle has the following differences:
 1. Action is a PL/SQL statement.
 2. New/old tuples referenced automatically.
 3. Strong constraints on trigger actions designed to make certain you can't fire off an infinite sequence of triggers.
- ◆ See on-line or-triggers.html document.

What Is a User-defined Function?

◆ Scalar Functions

- ◆ Similar to a built-in function

◆ Multi-Statement Table-valued Functions

- ◆ Content like a stored procedure
- ◆ Referenced like a view

◆ In-Line Table-valued Functions

- ◆ Similar to a view with parameters
- ◆ Returns a table as the result of single SELECT statement

Defining User-defined Functions

- ◆ Creating a User-defined Function
- ◆ Creating a Function with Schema Binding
- ◆ Setting Permissions for User-defined Functions
- ◆ Altering and Dropping User-defined Functions

Creating a User-defined Function

■ Creating a Function

```
USE Northwind
CREATE FUNCTION fn_NewRegion
    (@myinput nvarchar(30))
    RETURNS nvarchar(30)
BEGIN
    IF @myinput IS NULL
        SET @myinput = 'Not Applicable'
    RETURN @myinput
END
```

■ Restrictions on Functions

Creating a Function with Schema Binding

- ◆ Referenced User-defined Functions and Views Are Also Schema Bound
- ◆ Objects Are Not Referenced with a Two-Part Name
- ◆ Function and Objects Are All in the Same Database
- ◆ Have Reference Permission on Required Objects

Setting Permissions for User-defined Functions

- ◆ Need CREATE FUNCTION Permission
- ◆ Need EXECUTE Permission
- ◆ Need REFERENCE Permission on Cited Tables, Views, or Functions
- ◆ Must Own the Function to Use in CREATE or ALTER TABLE Statement

Altering and Dropping User-defined Functions

◆ Altering Functions

```
ALTER FUNCTION dbo.fn_NewRegion  
    <New function content>
```

- ◆ Retains assigned permissions
- ◆ Causes the new function definition to replace existing definition

◆ Dropping Functions

```
DROP FUNCTION dbo.fn_NewRegion
```

Examples of User-defined Functions

- ◆ Using a Scalar User-defined Function
- ◆ Example of a Scalar User-defined Function
- ◆ Using a Multi-Statement Table-valued Function
- ◆ Example of a Multi-Statement Table-valued Function
- ◆ Using an In-Line Table-valued Function
- ◆ Example of an In-Line Table-valued Function

Using a Scalar User-defined Function

- ◆ RETURNS Clause Specifies Data Type
- ◆ Function Is Defined Within a BEGIN and END Block
- ◆ Return Type Is Any Data Type Except text, ntext, image, cursor, or timestamp

Example of a Scalar User-defined Function

■ Creating the Function

```
CREATE FUNCTION fn_DateFormat
    (@indate datetime, @separator char(1))
    RETURNS Nchar(20)
    AS
    BEGIN
        RETURN
        CONVERT(Nvarchar(20), datepart(mm,@indate))
        + @separator
        + CONVERT(Nvarchar(20), datepart(dd, @indate))
        + @separator
        + CONVERT(Nvarchar(20), datepart(yy, @indate))
    END
```

■ Calling the Function

```
SELECT dbo.fn_DateFormat(GETDATE(), ':')
```

Using a Multi-Statement Table-valued Function

- ◆ BEGIN and END Enclose Multiple Statements
- ◆ RETURNS Clause Specifies table Data Type
- ◆ RETURNS Clause Names and Defines the Table

Example of a Multi-Statement Table-valued Function

◆ Creating the Function

```
CREATE FUNCTION fn_Employees (@length nvarchar(9))
RETURNS @fn_Employees table
    (EmployeeID int PRIMARY KEY NOT NULL,
     [Employee Name] nvarchar(61) NOT NULL)
AS
BEGIN
    IF @length = 'ShortName'
        INSERT @fn_Employees SELECT EmployeeID, LastName
        FROM Employees
    ELSE IF @length = 'LongName'
        INSERT @fn_Employees SELECT EmployeeID,
        (FirstName + ' ' + LastName) FROM Employees
RETURN
END
```

◆ Calling the Function

```
SELECT * FROM dbo.fn_Employees('LongName')
Or
SELECT * FROM dbo.fn_Employees('ShortName')
```

Using an In-Line Table-valued Function

- ◆ Content of the Function Is a SELECT Statement
- ◆ Do Not Use BEGIN and END
- ◆ RETURN Specifies table as the Data Type
- ◆ Format Is Defined by the Result Set

Example of an In-Line Table-valued Function

■ Creating the Function

```
CREATE FUNCTION fn_CustomerNamesInRegion
    ( @RegionParameter nvarchar(30) )
RETURNS table
AS
RETURN (
    SELECT CustomerID, CompanyName
    FROM Northwind.dbo.Customers
    WHERE Region = @RegionParameter
)
```

■ Calling the Function Using a Parameter

```
SELECT * FROM fn_CustomerNamesInRegion(N'WA')
```

SQLPlus

- ◆ In addition to stored procedures, one can write a PL/SQL statement that looks like the body of a procedure, but is executed once, like any SQL statement typed to the generic interface.
 - ◆ Oracle calls the generic interface “sqlplus.”
 - ◆ PL/SQL is really the “plus.”

Form of PL/SQL Statements

DECLARE

<declarations>

BEGIN

<statements>

END;

.

run

◆ The DECLARE section is optional.

Form of PL/SQL Procedure

CREATE OR REPLACE PROCEDURE

<name> (<arguments>) **AS** ← Notice AS
needed here

<optional declarations>

BEGIN

<PL/SQL statements>

END;

•
run

← Needed to store
procedure in database;
does not really run it.

PL/SQL Declarations and Assignments

- ◆ The word DECLARE does not appear in front of each local declaration.
 - ◆ Just use the variable name and its type.
- ◆ There is no word SET in assignments, and := is used in place of =.
 - ◆ **Example:** `x := y;`

PL/SQL Procedure Parameters

- ◆ There are several differences in the forms of PL/SQL argument or local-variable declarations, compared with the SQL/PSM standard:
 1. Order is name-mode-type, not mode-name-type.
 2. INOUT is replaced by IN OUT in PL/SQL.
 3. Several new types.

PL/SQL Types

- ◆ In addition to the SQL types, NUMBER can be used to mean INT or REAL, as appropriate.
- ◆ You can refer to the type of attribute x of relation R by $R.x\%TYPE$.
 - ◆ Useful to avoid type mismatches.
 - ◆ Also, $R\%ROWTYPE$ is a tuple whose components have the types of R 's attributes.

Example:JoeMenu

- ◆ Recall the procedure **JoeMenu(b,p)** that adds beer b at price p to the beers sold by Joe (in relation Sells).
- ◆ Here is the PL/SQL version.

Procedure JoeMenu in PL/SQL

```
CREATE OR REPLACE PROCEDURE JoeMenu (  
  b IN Sells.beer%TYPE,  
  p IN Sells.price%TYPE  
) AS  
  BEGIN  
    INSERT INTO Sells  
    VALUES ('Joe''s Bar', b, p);  
  END;
```

Notice these types
will be suitable
for the intended
uses of *b* and *p*.

```
·  
run
```

PL/SQL Branching Statements

- ◆ Like IF ... in SQL/PSM, but:
- ◆ Use ELSIF in place of ELSEIF.
- ◆ Viz.: IF ... THEN ... ELSIF ... THEN ...
ELSIF ... THEN ... ELSE ... END IF;

PL/SQL Loops

- ◆ LOOP ... END LOOP as in SQL/PSM.
- ◆ Instead of LEAVE ... , PL/SQL uses
EXIT WHEN <condition>
- ◆ And when the condition is that cursor *c* has found no tuple, we can write *c*%NOTFOUND as the condition.

PL/SQL Cursors

- ◆ The form of a PL/SQL cursor declaration is:

CURSOR <name> IS <query>;

- ◆ To fetch from cursor c, say:

FETCH c INTO <variable(s)>;

Example: JoeGouge() in PL/SQL

- ◆ Recall **JoeGouge()** sends a cursor through the Joe's-Bar portion of Sells, and raises by \$1 the price of each beer Joe's Bar sells, if that price was initially under \$3.

Example: JoeGouge() Declarations

```
CREATE OR REPLACE PROCEDURE
    JoeGouge() AS
theBeer Sells.beer%TYPE;
thePrice Sells.price%TYPE;
CURSOR c IS
    SELECT beer, price FROM Sells
    WHERE bar = 'Joe''s Bar';
```

Example: JoeGouge() Body

```
BEGIN
  OPEN c;
  LOOP
    FETCH c INTO theBeer, thePrice;
    EXIT WHEN c%NOTFOUND;
    IF thePrice < 3.00 THEN
      UPDATE Sells SET price = thePrice + 1.00;
      WHERE bar = 'Joe"s Bar" AND beer = theBeer;
    END IF;
  END LOOP;
  CLOSE c;
END;
```

How PL/SQL breaks a cursor loop

Note this is a SET clause in an UPDATE, not an assignment. PL/SQL uses := for assignments.

Tuple-Valued Variables

- ◆ PL/SQL allows a variable x to have a tuple type.
- ◆ $x \text{ R\%ROWTYPE}$ gives x the type of R 's tuples.
- ◆ R could be either a relation or a cursor.
- ◆ $x.a$ gives the value of the component for attribute a in the tuple x .

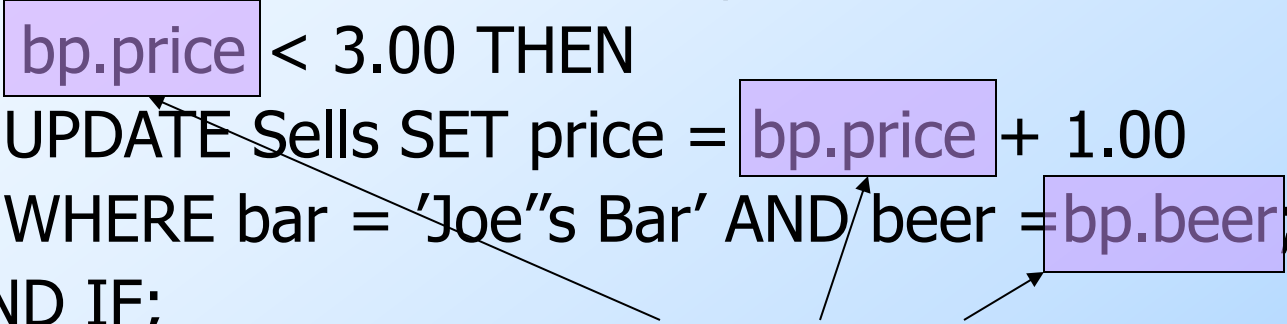
Example: Tuple Type

- ◆ Repeat of JoeGouge() declarations with variable *bp* of type beer-price pairs.

```
CREATE OR REPLACE PROCEDURE
    JoeGouge () AS
    CURSOR c IS
    SELECT beer, price FROM Sells
    WHERE bar = 'Joe''s Bar';
    bp c%ROWTYPE;
```

JoeGouge() Body Using *bp*

```
BEGIN
  OPEN c;
  LOOP
    FETCH c INTO bp;
    EXIT WHEN c%NOTFOUND;
    IF bp.price < 3.00 THEN
      UPDATE Sells SET price = bp.price + 1.00
      WHERE bar = 'Joe"s Bar' AND beer = bp.beer;
    END IF;
  END LOOP;
  CLOSE c;
END;
```



Components of bp are obtained with a dot and the attribute name

Embedded SQL

- ◆ **Key idea:** A preprocessor turns SQL statements into procedure calls that fit with the surrounding host-language code.
- ◆ All embedded SQL statements begin with EXEC SQL, so the preprocessor can find them easily.

Shared Variables

- ◆ To connect SQL and the host-language program, the two parts must share some variables.
- ◆ Declarations of shared variables are bracketed by:

 `EXEC SQL BEGIN DECLARE SECTION;`

`<host-language declarations>`

 `EXEC SQL END DECLARE SECTION;`

Use of Shared Variables

- ◆ In SQL, the shared variables must be preceded by a colon.
 - ◆ They may be used as constants provided by the host-language program.
 - ◆ They may get values from SQL statements and pass those values to the host-language program.
- ◆ In the host language, shared variables behave like any other variable.

Example: Looking Up Prices

- ◆ We'll use C with embedded SQL to sketch the important parts of a function that obtains a beer and a bar, and looks up the price of that beer at that bar.
- ◆ Assumes database has our usual `Sells(bar, beer, price)` relation.

Example: C Plus SQL

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char theBar[21], theBeer[21];
```

```
float thePrice;
```

```
EXEC SQL END DECLARE SECTION;
```

```
/* obtain values for theBar and theBeer */
```

```
EXEC SQL SELECT price INTO :thePrice
```

```
FROM Sells
```

```
WHERE bar = :theBar AND beer = :theBeer;
```

```
/* do something with thePrice */
```

Note 21-char
arrays needed
for 20 chars +
endmarker

SELECT-INTO
as in PSM 75

Embedded Queries

- ◆ Embedded SQL has the same limitations as PSM regarding queries:
 - ◆ SELECT-INTO for a query guaranteed to produce a single tuple.
 - ◆ Otherwise, you have to use a cursor.
 - Small syntactic differences, but the key ideas are the same.

Cursor Statements

- ◆ Declare a cursor *c* with:

```
EXEC SQL DECLARE c CURSOR FOR <query>;
```

- ◆ Open and close cursor *c* with:

```
EXEC SQL OPEN CURSOR c;
```

```
EXEC SQL CLOSE CURSOR c;
```

- ◆ Fetch from *c* by:

```
EXEC SQL FETCH c INTO <variable(s)>;
```

- ◆ Macro NOT FOUND is true if and only if the FETCH fails to find a tuple.


Example: Print Joe's Menu

- ◆ Let's write C + SQL to print Joe's menu
 - the list of beer-price pairs that we find in `Sells(bar, beer, price)` with `bar = Joe's Bar`.
- ◆ A cursor will visit each Sells tuple that has `bar = Joe's Bar`.

Example: Declarations

```
EXEC SQL BEGIN DECLARE SECTION;  
    char theBeer[21]; float thePrice;  
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL DECLARE c CURSOR FOR  
    SELECT beer, price FROM Sells  
    WHERE bar = 'Joe''s Bar';
```



The cursor declaration goes
outside the declare-section

Example: Executable Part

```
EXEC SQL OPEN CURSOR c;
```

```
while(1){
```

```
    EXEC SQL FETCH c
```

```
        INTO :theBeer, :thePrice;
```

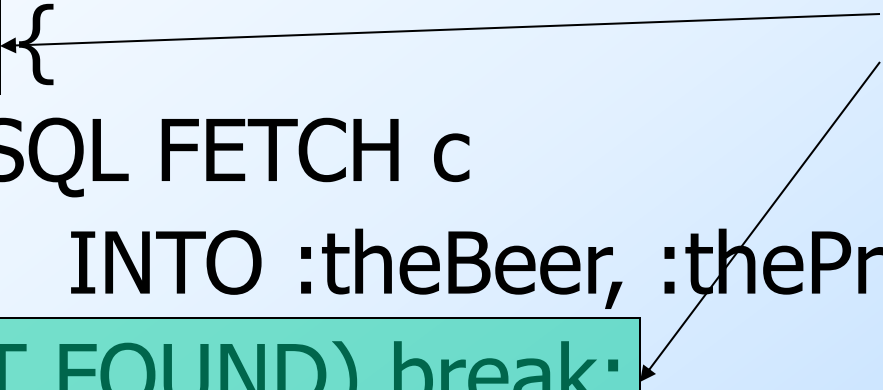
```
    if (NOT FOUND) break;
```

```
    /* format and print theBeer and thePrice */
```

```
}
```

```
EXEC SQL CLOSE CURSOR c;
```

The C style
of breaking
loops



Need for Dynamic SQL

- ◆ Most applications use specific queries and modification statements to interact with the database.
 - ◆ The DBMS compiles EXEC SQL ... statements into specific procedure calls and produces an ordinary host-language program that uses a library.
- ◆ What about sqlplus, which doesn't know what it needs to do until it runs?

Dynamic SQL

- ◆ Preparing a query:

```
EXEC SQL PREPARE <query-name>  
        FROM <text of the query>;
```

- ◆ Executing a query:

```
EXEC SQL EXECUTE <query-name>;
```

- ◆ “Prepare” = optimize query.
- ◆ Prepare once, execute many times.

Example: A Generic Interface

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    char query[MAX_LENGTH];
```

```
EXEC SQL END DECLARE SECTION;
```

```
while(1) {
```

```
    /* issue SQL> prompt */
```

```
    /* read user's query into array query */
```

```
    EXEC SQL PREPARE q FROM :query;
```

```
    EXEC SQL EXECUTE q;
```

```
}
```

q is an SQL variable
representing the optimized
form of whatever statement
is typed into :query

Execute-Immediate

- ◆ If we are only going to execute the query once, we can combine the PREPARE and EXECUTE steps into one.

◆ Use:

```
EXEC SQL EXECUTE IMMEDIATE <text>;
```

Example: Generic Interface Again

```
EXEC SQL BEGIN DECLARE SECTION;
    char query[MAX_LENGTH];
EXEC SQL END DECLARE SECTION;
while(1) {
    /* issue SQL> prompt */
    /* read user's query into array
    query */
    EXEC SQL EXECUTE IMMEDIATE :query;
}
```