

# Chapter 15: View Serializability

# View Serializability

Conflict equivalent

View equivalent

Conflict serializable

View serializable

# Motivating example

## Schedule Q

T<sub>1</sub>

T<sub>2</sub>

T<sub>3</sub>

Read(A)

Write(A)


Write(A)

Write(A)

Same as

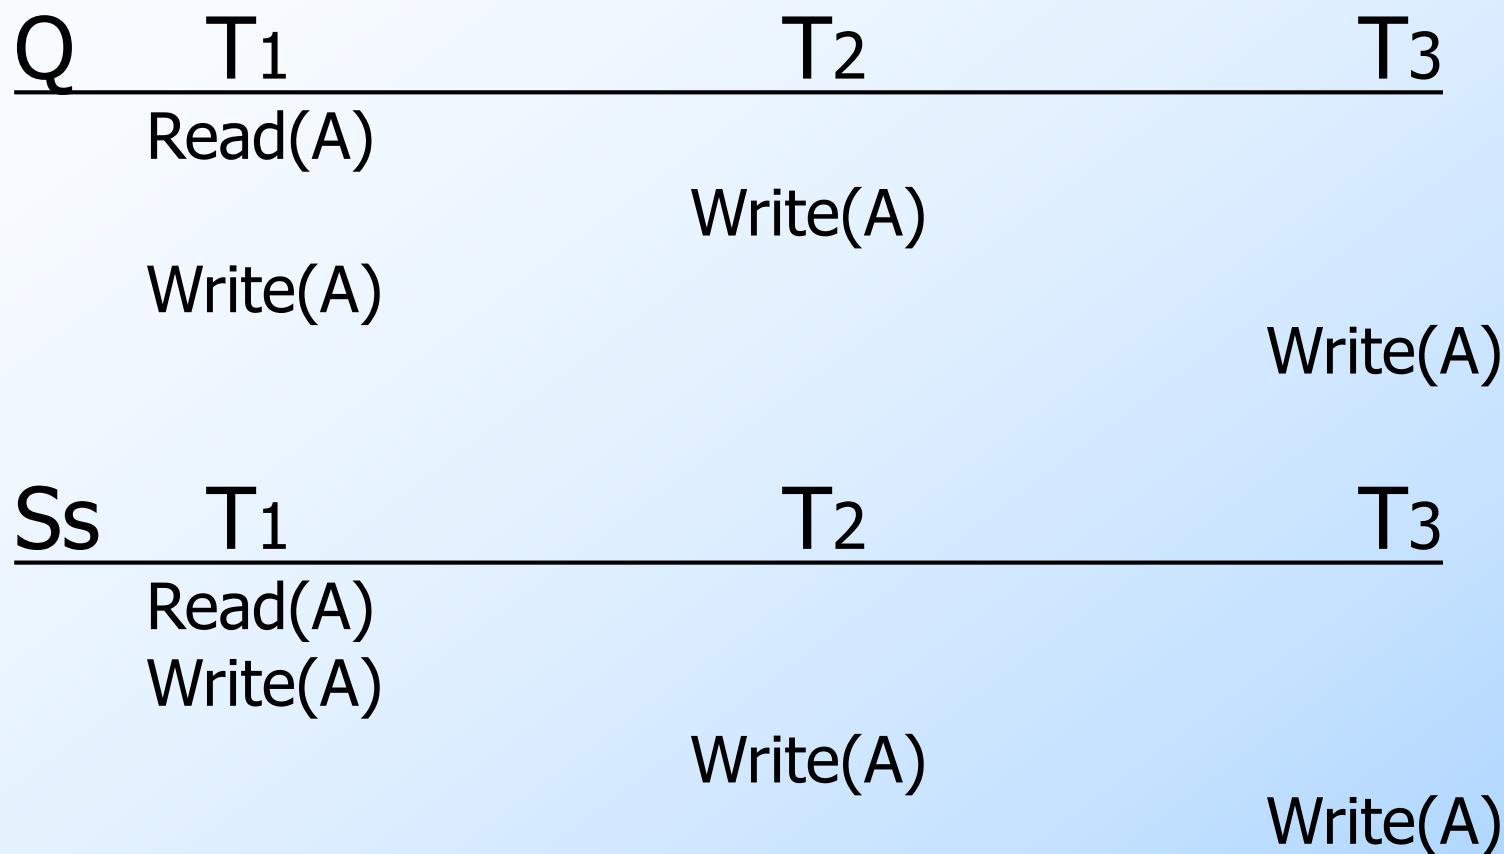
$Q = r_1(A) \ w_2(A) \ w_1(A) \ w_3(A)$

$P(Q): T_1 \longleftrightarrow T_2$



Not conflict serializable!

But now compare Q to Ss, a serial schedule:



- ◆  $T_1$  reads same thing in  $Q$ ,  $S_s$
- ◆  $T_2, T_3$  read something (nothing?)
- ◆ After  $Q$  or  $S_s$ , DB is left in same state

So what is wrong with  $Q$ ?

Definition Schedules  $S_1, S_2$  are  
View Equivalent if:

- (1) If in  $S_1$ :  $w_j(A) \Rightarrow r_i(A)$   
then in  $S_2$ :  $w_j(A) \Rightarrow r_i(A)$
- (2) If in  $S_1$ :  $r_i(A)$  reads initial DB value,  
then in  $S_2$ :  $r_i(A)$  also reads initial DB  
value
- (3) If in  $S_1$ :  $T_i$  does last write on  $A$ ,  
then in  $S_2$ :  $T_i$  also does last write on  $A$

$\Rightarrow$  means "reads  
value produced"

## Definition

Schedule  $S_1$  is View Serializable if it is view equivalent to some serial schedule



View  
Serializable

← ? →

Conflict  
Serializable

- View Serializable  $\not\Rightarrow$  Conflict Serializable  
e.g., See Schedule Q
- Conflict Serializable  $\stackrel{?}{\Rightarrow}$  View Serializable

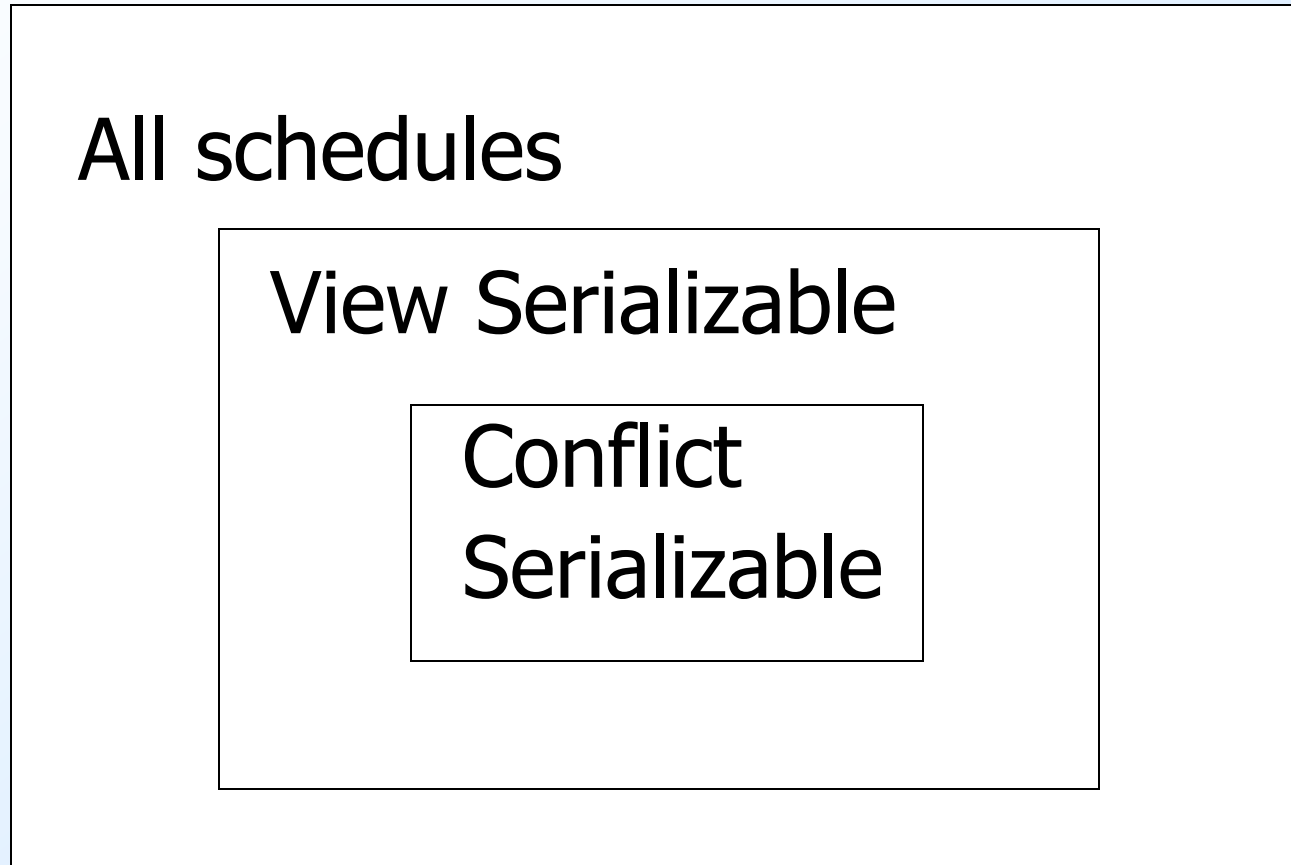
## Lemma

Conflict Serializable  $\Rightarrow$  View Serializable

### Proof:

— Swapping non-conflicting actions does not change what transactions read nor final DB state

# Venn Diagram



Note: All view serializable schedules that are not conflict serializable, involve useless write

$$S = W_2(A) \xleftrightarrow{\text{no reads}} W_3(A) \dots$$

FALSE: Counterexample (Sorav Bansal):  
 $w_3(Y) \ r_2(Y) \ w_1(X) \ r_2(X) \ w_3(X) \ r_4(X) \ w_5(X)$

# How do we test for view-serializability?


P(S) not good enough...  
(see schedule Q)

- ◆ One problem: some swaps involving conflicting actions are OK... e.g.:

$S = \dots w_2(A) \dots r_1(A) \dots w_3(A) \dots w_4(A)$

this action can move  
if this write exists - - - - -

## ◆ Another problem: useless writes



$S = \dots W_2(A) \dots W_1(A) \dots$

no A reads

# To check if S is View Serializable

(1) Add final transaction  $T_f$  that reads all DB

(eliminates condition 3 of V-S definition)

E.g.:  $S = \dots W_1(A) \dots W_2(A) \dots rf(A)$

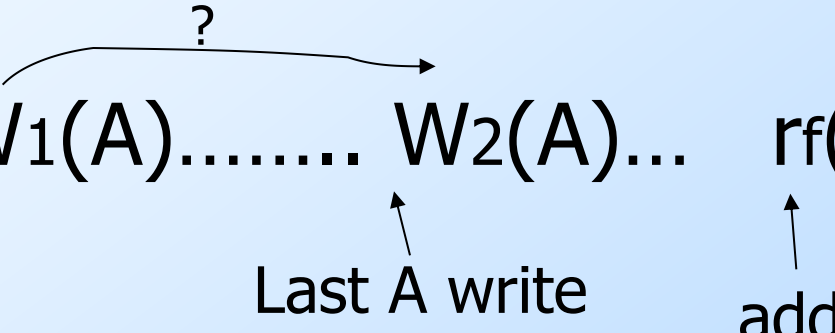


Diagram illustrating a sequence of transactions  $S$  with annotations:

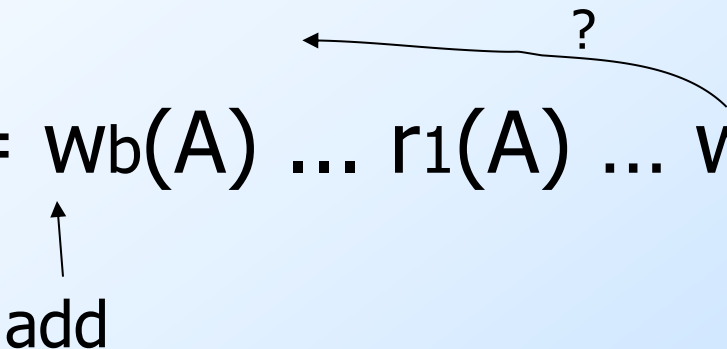
- A curved arrow labeled  $?$  points from  $W_1(A)$  to  $W_2(A)$ .
- An arrow labeled "Last A write" points to  $W_2(A)$ .
- An arrow labeled "add" points to  $rf(A)$ .



(2) Add initial transaction  $T_b$   
that writes all DB

(eliminates condition 2 of V-S definition)

E.g.:  $S = w_b(A) \dots r_1(A) \dots w_2(A) \dots$



The diagram includes two annotations: a vertical arrow pointing upwards from the word 'add' to the transaction  $w_b(A)$ , and a curved arrow starting from a question mark '?' and pointing to the transaction  $w_2(A)$ .

0

(3) Create labeled precedence graph of  $S$ :

(3a) If  $w_i(A) \Rightarrow r_j(A)$  in  $S$ , add  $T_i \rightarrow T_j$

(3b) For each  $w_i(A) \Rightarrow r_j(A)$  do

consider each  $w_k(A)$ :  $[T_k \neq T_b]$

- If  $T_i \neq T_b \wedge T_j \neq T_f$  then insert

$$\begin{cases} T_k \xrightarrow{p} T_i \\ T_j \xrightarrow{p} T_k \end{cases} \quad \text{some new } p$$

- If  $T_i = T_b \wedge T_j \neq T_f$  then insert

$$T_j \xrightarrow{0} T_k$$

- If  $T_i \neq T_b \wedge T_j = T_f$  then insert

$$T_k \xrightarrow{0} T_i$$


(4) Check if  $LP(S)$  is “acyclic” (if so,  $S$  is  $V-S$ )

- For each pair of “ $p$ ” arcs ( $p \neq 0$ ),  
choose one

Example: check if Q is V-S:

$Q = r_1(A) \ w_2(A) \ w_1(A) \ w_3(A)$


$Q' = w_b(A) \Rightarrow r_1(A) \ w_2(A) \ w_1(A) \ w_3(A) \Rightarrow r_f(A)$



LP(S) acyclic!!  
S is V-S

Another example:

$Z = w_b(A) \Rightarrow r_1(A) \quad w_2(A) \Rightarrow r_3(A) \quad w_1(A) \quad w_3(A) \Rightarrow r_f(A)$



do not pick  
this one of "1" pair

LP(Z) acyclic, so Z is V-S  
(equivalent to  $T_b \ T_1 \ T_2 \ T_3 \ T_f$ )

$$S_s = w_b(A) \underbrace{r_1(A)w_1(A)}_{T_1} \underbrace{w_2(A)r_3(A)}_{T_2} \underbrace{w_3(A)r_f(A)}_{T_3}$$

$Z + S_s$  indeed do same thing

- ◆ Checking view serializability is expensive
- ◆ Still, V-S useful in some cases...




## Example on useless transactions:

$S = w_1(A) \ r_2(A) \ w_2(B) \ r_1(B) \ w_3(A) \ w_3(B)$

$S' =$

$T_b \ w_1(A) \Rightarrow r_2(A) \ w_2(B) \Rightarrow r_1(B) \ w_3(A) \ w_3(B) \Rightarrow T_f$



- ◆ If we only care about final state  
remove  $T_1, T_2$ ; i.e., remove useless  
transactions
- ◆ If we care what  $T_1, T_2$  read (view  
equivalence), then do not remove  
useless transactions

◆ If all transactions read what they write,  
(I.e.,  $T_j = \dots R_j(A) \dots W_j(A) \dots$ ) then  
view serializability = conf. serializability

[Another way of saying: blind writes  
appear in any view-serializable schedule  
that is not conflict serializable]

Proof(?): say  $S_1$  is view-ser. and no blind writes.  $S_1$  V-equiv to  $S_s$ , serial schedule.

(1) Goal: Show that

$$T_1 \rightarrow T_2 \text{ in } P(S_1) \Rightarrow T_1 <_{ss} T_2$$

(2) Assume  $T_1 \rightarrow T_2$

if  $S_1 = \dots w_1(A) \dots r_2(A) \dots$

(direct read) clearly  $T_1 <_{ss} T_2$

if  $S_1 = \dots w_1(A) \dots r_3(A) w_3(A) \dots r_2(A) \dots$

also  $T_1 <_{ss} T_2$

if  $S_1 = \dots r_1(A) r_3(A) \dots w_1(A) \dots w_3(A) \dots r_2(A)$

not possible:  $T_1, T_3$  not

serializable

Other cases similar...

## Implications:

If no blind writes, view-ser  $\Leftrightarrow$  conf-ser

$P(S)$  acyclic  $\Rightarrow$  all transactions read the same as in a serial schedule